

Fairness-Enhancing Classification Methods for Non-Binary Sensitive Features

How to Fairly Detect Leakages in Water Distribution Systems

Janine Strotherm, Inaam Ashraf and Barbara Hammer

Center for Cognitive Interaction Technology, Bielefeld University, Bielefeld, Germany
{jstrotherm,mashraf,bhammer}@techfak.uni-bielefeld.de

Corresponding Author: Janine Strotherm

Abstract

Especially if AI-supported decisions affect the society, the fairness of such AI-based methodologies constitutes an important area of research. In this contribution, we investigate the applications of AI to the socioeconomically relevant infrastructure of water distribution systems (WDSs). We propose an appropriate definition of protected groups in WDSs and generalized definitions of group fairness that provably coincide with existing definitions in their specific settings. We demonstrate that typical methods for the detection of leakages in WDSs are unfair in this sense. Further, we thus propose a general fairness-enhancing framework as an extension of the specific leakage detection pipeline, but also for an arbitrary learning scheme, to increase the fairness of the AI-based algorithm. Finally, we evaluate and compare several specific instantiations of this framework on a toy and on a realistic WDS to show their utility.

Keywords: Fairness · Fair Machine Learning · Independence · Separation · Disparate Impact · Demographic Parity · Equal Opportunity · Equalized Odds · Leakage Detection · Water Distribution Networks

1 Introduction

Due to the increasing usage of artificial intelligence (AI)-based decision making systems in socially relevant fields of application the question of *fair decision making* gained much importance in recent years (cf. [4], [15]). Fairness in is hereby related to the several (protected) groups or individuals which are affected by the algorithmic decision making and characterized by *sensitive features* such as gender or ethnicity. Most algorithms on which these tools are based on rely on data which can be biased with respect to the questions of fairness without intention, resulting in skewed models. Also the algorithm itself can discriminate protected groups or individuals without explicitly aiming to do so due to an undesirable algorithmic bias (cf. [26, 29]). This gives rise to the question on how to define fairness and how to mitigate unfairness in case it occurs in the context of machine learning (ML), i.e., in the context of data-driven algorithms.

Background: Fairness definitions

Several definitions of fairness as well as approaches to achieve these fairness standards have been theoretically discussed and tested in practice (cf. [7, 12, 14, 26, 29]). From a legal perspective, one distinguishes between *disparate treatment* and *disparate impact* (cf. [7]). While disparate treatment occurs whenever a group or an individual is intentionally treated differently because of their membership in a protected group, disparate impact is a consequence of indirect discrimination happening despite “seemingly neutral policy” (cf. [29]).

From a scientific viewpoint, the variety of fairness notions is much larger where many popular approaches focus mainly on (binary) classification tasks (cf. [12, 26, 29]). Different definitions can be grouped into the concepts of *group fairness*, *individual fairness*, *causal fairness* and *dynamic fairness*: Group fairness aims at treating different groups equally while individual fairness aims at treating similar individuals similarly. Causal fairness examines the extent to which the sensitive feature, such as gender or ethnicity, influences the prediction of a model and dynamic fairness examines the long-term effects of (supposedly) fair decisions (cf. [36]).

The fairness notions that we will discuss in this work belong to the former concept of group fairness. Here, most works focus on fairness definitions with respect to a single binary sensitive feature that splits the underlying population in a *discriminated* and a *privileged* group (cf. [16, 19, 21, 22, 26, 29, 32, 42, 43]). There is few work on fairness definitions based on the independence assumption of the model’s prediction and a single non-binary sensitive feature, however, there is no rigorous theory on how this assumption translates to generalized fairness notions as necessary and sufficient conditions of this independence assumption and their relation to the binary case (cf. [1, 12]). We will build on this point.

Background: Fairness methods

Besides the definition of fairness, the problem arises how to enhance fairness in well known ML methods while maintaining a reasonable overall performance of the model. Approaches can hereby be grouped into three categories: Depending on when in the training pipeline the model is enhanced with respect to fairness, we speak about pre-, in- or post-processing techniques (cf. [7, 26, 29]).

Pre-processing methods usually modify the training data which is fed to the training algorithm. For example, Kamiran et al. [22] use a resampling technique by removing unpreferred samples, i.e., positive outcomes in the privileged group and negative outcomes in the discriminated group, and duplicating preferred samples, i.e., positive outcomes in the discriminated group and negative outcomes in the privileged group, that lie close to the decision boundary of a binary classifier. In another work, they modify the training data by changing the labels of training samples that lie close to the decision boundary of a binary classifier such that negative outcomes in the privileged group and positive outcomes in the discriminated group appear more often (cf. [21]). While these methods aim at putting more emphasis on the discriminated group and less emphasis on the privileged group, Feldman et al. [16] modify the non-sensitive features of training samples such that it is not able to predict the sensitive feature from the non-sensitives. This reduces the chance that the model’s predictions, which are based on the non-sensitive features, are correlated to the sensitive feature.

In contrast, post-processing methods modify the model after the training. For example, Pleiss et al. [30] modify a pre-trained model by randomly changing some outputs of a binary classifier on the group on which the classifier performs better to ensure equal performance over all groups. As another example, Hardt et al. [19] retrain a pre-trained model by optimizing a loss between the new and the pre-trained binary classifier while

satisfying fairness-constraints. Another simple approach is to use group-specific thresholds for a threshold-based classifier (cf. [13]).

Finally, in-process methods modify the (original) training algorithm. A common way to do so is by adding fairness-constraints (cf. [1, 2, 11, 24, 27, 43, 42]) or a fairness-regularization-term (cf. [3, 9, 29]) to the to be optimized loss function. Next to classification, also regression tasks usually fall into this category (cf. [2, 3, 9, 11, 24, 27]).

The methods presented in this work are also in-process methods and are extensions of our work [35], published in Springer’s Lecture Notes in Computer Science. Both of these works are based on the methods of Zafar et al. [43], but adapted to more generalized settings, as we will elaborate in the contributions paragraph.

Background: Fairness in water distribution systems (WDSs) The question of fairness becomes especially relevant when the decisions of a ML model take impact on socioeconomic infrastructure, such as WDSs. To the best of our knowledge, our previous work [35] has been the first approach to introduce fairness within this domain. In this work, we address the important problem of leakage detection in WDSs and investigate in how far typical models treat different groups of consumers of the WDS (in)equally, and we will extend these considerations in this work.

Contributions Our approaches to improve group fairness in such a domain of high social and ethical relevance are based on the idea of considering the locality in the WDS as a sensitive feature. Considering the empirical covariance between the sensitive feature(s) and the model’s prediction as a proxy for the fairness measure, similar to Zafar et al. [43], but also the generalized fairness notions directly, are the base of all our proposed methods. The advantage of our fairness-enhancing algorithms is that they can handle even multiple non-binary sensitive features and satisfy both the concept of disparate treatment *and* disparate impact simultaneously, which is an asset towards most fairness-enhancing algorithms (cf. [29, 43]). In more detail, our contributions – also in view of what this extension offers compared to our previous work [35] – are as follows:

- We propose group fairness definitions even for multiple non-binary sensitive features and which are generalizations of well-known corresponding fairness notions in the common setting of a binary classifier and a single binary sensitive feature.
- As an extension to our previous work, we provide details on the mathematical concept of independence, derive easily-to-test independence criteria and leverage these in order to derive those generalized group fairness definitions.
- As an extension to our previous work, we prove that those coincide with the before-mentioned well-known corresponding fairness notions.
- We introduce a common leakage detection pipeline and propose a suitable definition of sensitive features and group fairness in the context of leakage detection in WDSs, with more detail in this work compared to our previous work.
- We present specific instantiations of that pipeline and show that common leakage detection methods do not obey those fairness criteria, with one more specific instantiation (based on the more powerful graph convolutional network (GCN) based virtual sensors instead of linear regression based virtual sensors) and with more detail in this work compared to our previous work.

- We introduce a fairness-enhancing leakage detection framework as an extension of the common leakage detection pipeline, with more detail in this work compared to our previous work.
- We present specific instantiations of this framework, among others by modifying the ideas of Zafar et al. [43] to any (ensemble) classification model instead of a convex margin-based binary classifier, to propose several fairness-enhancing methods, with more specific instantiations, among others based on the ideas in our previous work made on possible modifications of our methodologies.
- We provide an empirical evaluation of our proposed methods.
- As an extension to our previous work, next to the application of those methods to the toy WDS *Hanoi*, we investigate on the application to the more complex and realistic WDS *L-Town*.

Structure of the work The rest of this work is structured as followed: In section 2, we introduce definitions of group fairness for multiple non-binary sensitive features, giving the mathematical background for the derivation of such generalized definitions and how they are connected to already existing definitions. Afterwards, in section 3, we present a standard methodology to detect leakages in WDSs, introduce the meaning of sensitive features in this context and investigate whether the resulting model makes fair decisions with respect to the previously defined notions of fairness. Consecutively, in section 4, we propose and evaluate several adaptations of this methodology that enhance fairness and provide empirical evidence for our theoretical findings regarding the equivalence of different fairness notions in this specific domain of application. Finally, our findings are summarized and discussed in section 5.

2 Group Fairness In Machine Learning

On an abstract level, the concept of group fairness is based on the mathematical concept of (conditional) independence of two random variables (cf. [7, 12]). Therefore, in this section, we will first investigate on this concept of independence in general (subsection 2.1). Consecutively, we will introduce the mathematical notation required to define a ML task and group fairness of such based on this general concept of independence (subsection 2.2) to be able to derive group fairness definitions in generalized ML tasks, which coincide with well-known definitions in more specific settings (subsection 2.3).

2.1 Independence of Two Random Variables

As it is the main mathematical concept to characterize different notions of group fairness, for the sake of convenience, we recapitulate the concept of independence of two random variables in this subsection. Moreover, we target an easy necessary and sufficient condition for this concept which is particularly simple to apply and test in the context of fairness of ML models. Hence, we derive an equivalent formulation, lemma 2.2, which can be tested on canonical subsets of the full σ -fields.

For the rest of this subsection, let $(\Omega, \mathcal{F}, \mathbb{P})$ be a probability space, $(\mathcal{X}, \mathcal{F}_{\mathcal{X}})$, $(\mathcal{Y}, \mathcal{F}_{\mathcal{Y}})$ measurable spaces and $X : \Omega \rightarrow \mathcal{X}$, $Y : \Omega \rightarrow \mathcal{Y}$ random variables¹.

Definition 2.1 (Independence of two random variables (cf. [8])).

X and Y are independent with respect to the probability measure \mathbb{P} , if the σ -fields² $\sigma(X) := \{X^{-1}(A) \mid A \in \mathcal{F}_{\mathcal{X}}\} \subset \mathcal{F}$ and $\sigma(Y) := \{Y^{-1}(B) \mid B \in \mathcal{F}_{\mathcal{Y}}\} \subset \mathcal{F}$ generated by these random variables are independent with respect to \mathbb{P} .

Based on that, in appendix A.2, we derive general necessary and sufficient conditions for independence of two random variables. In the context of fairness of ML models, we are usually interested into a more specific setting, namely into the independence of two discrete random variables.

Lemma 2.2 (Independence of two discrete random variables).

Assume that X and Y are discrete, i.e., that $\mathcal{X} = \{x_1, \dots, x_{K_x}\}$ and $\mathcal{Y} = \{y_1, \dots, y_{K_y}\}$ holds. Then X and Y are independent with respect to \mathbb{P} if and only if (iff)

$$\mathbb{P}(X = x \mid Y = y_{k_1}) = \mathbb{P}(X = x \mid Y = y_{k_2}) \quad (2.1)$$

holds for all $x \in \mathcal{X}$ and $y_{k_1}, y_{k_2} \in \mathcal{Y}$ for which $\mathbb{P}(Y = y_{k_1}), \mathbb{P}(Y = y_{k_2}) > 0$ holds.

Proof. $\mathcal{E}_{\mathcal{X}} = \{\emptyset, \{x_1\}, \dots, \{x_{K_x}\}\}$ and $\mathcal{E}_{\mathcal{Y}} = \{\emptyset, \{y_1\}, \dots, \{y_{K_y}\}\}$ are \cap -stable generators of the σ -fields $\mathcal{F}_{\mathcal{X}} = P(\mathcal{X}) = \sigma(\mathcal{E}_{\mathcal{X}})$ and $\mathcal{F}_{\mathcal{Y}} = P(\mathcal{Y}) = \sigma(\mathcal{E}_{\mathcal{Y}})$, respectively. Therefore, by remark A.9 and lemma A.12 (we can replace the σ -fields $\mathcal{F}_{\mathcal{X}}$ and $\mathcal{F}_{\mathcal{Y}}$ in lemma A.12 by their generators $\mathcal{E}_{\mathcal{X}}$ and $\mathcal{E}_{\mathcal{Y}}$, respectively), X and Y are independent iff

$$\mathbb{P}(X \in \{x\} \mid Y \in \{y_{k_1}\}) = \mathbb{P}(X \in \{x\} \mid Y \in \{y_{k_2}\})$$

holds for all $x \in \mathcal{X} = \{x_1, \dots, x_{K_x}\}$ and all $y_{k_1}, y_{k_2} \in \mathcal{Y} = \{y_1, \dots, y_{K_y}\}$ for which $\mathbb{P}(Y \in \{y_{k_1}\}), \mathbb{P}(Y \in \{y_{k_2}\}) > 0$ holds. Note that we can omit the cases $X \in \emptyset$ and $Y \in \emptyset$, as these are trivially fulfilled. \square

Lemma 2.2 guarantees independence of two discrete random variables by only testing one-elementary events. As the σ -fields $\mathcal{F}_{\mathcal{X}} = P(\mathcal{X})$ and $\mathcal{F}_{\mathcal{Y}} = P(\mathcal{Y})$, given by *all* subsets of \mathcal{X} and \mathcal{Y} , respectively, consist of many more non-trivial events, lemma 2.2 gives us a valuable necessary and sufficient condition for independence of two discrete random variables, which we will make use of in the setting of ML.

2.2 Mathematical Notation for Machine Learning

Our next goal is the mathematical definition of group fairness as a formalization of equal treatment of a ML model *independent of sensitive attributes*. Such attributes, also called *sensitive features*, provide information about the membership or non-membership of a protected group, such as gender or ethnicity, in relation to which the model should not exhibit any prejudice (cf. [26, 29]). A later goal of this work will be to find a reasonable meaning of sensitive information in the context of WDS.

¹The statements in this subsection do not only hold for random variables (i.e., $\mathcal{X} \subset \mathbb{R}$, $\mathcal{Y} \subset \mathbb{R}$), but also for random vectors (i.e., $\mathcal{X} \subset \mathbb{R}^{d_x}$, $\mathcal{Y} \subset \mathbb{R}^{d_y}$) and even for random elements (i.e., \mathcal{X} and \mathcal{Y} are arbitrary sets). We just use the term random *variable* as this is the more familiar term.

²Of course, it requires a proof to show that these set systems are σ -fields. We give more information on this topic in appendix A.1 and especially, lemma A.4.

To formalize this independence and derive easily-to-test notions of group fairness based on the previous subsection 2.1 in the next subsection 2.3, we need to introduce mathematical notation that allows to consider independence of two random variables in the context of ML. In such context, probabilities such as in subsection 2.1 appear for random variables such as the model's output $\hat{Y} : \Omega \rightarrow \mathcal{Y}$, the labels $Y : \Omega \rightarrow \mathcal{Y}$, the features $X : \Omega \rightarrow \mathcal{X}$ or, in case of *fair* ML, the sensitive features $S : \Omega \rightarrow \mathcal{S}$ with target space \mathcal{Y} , feature space \mathcal{X} and sensitive feature space \mathcal{S} .

In this work, we will consider a one-dimensional binary classification task and a single discrete but possibly non-binary sensitive feature³, i.e., the target space equals $\mathcal{Y} = \{0, 1\} \subset \mathbb{N}_0$ and the (finite) sensitive feature space equals $\mathcal{S} = \{s_1, \dots, s_K\} \subset \mathbb{N}_0$ with. Equipping each with the power set makes them a measurable space.

Example 2.3 (Getting an intuition on fair ML). The domain Ω could consist of criminals in the US and the model \hat{Y} could predict whether ($\hat{Y}(\omega) = 1$) or whether not ($\hat{Y}(\omega) = 0$) a criminal $\omega \in \Omega$ will be criminal again in the future. This prediction should be independent on their ethnicity $S(\omega) \in \mathbb{N}_0$ (cf. [4]).

The typical goal of ML is to learn the relation between the features X and the labels Y , i.e., either the distribution $\mathbb{P}(X, Y)^{-1} = \mathbb{P}((X, Y) \in \cdot)$ of (X, Y) (generative ML) or more often, the distribution $\mathbb{P}(Y \in \cdot \mid X = x)$ of Y , given X , (discriminative ML) for any $x \in \mathcal{X}$.⁴ However, as these distributions are usually unknown, we use training data, i.e., samples⁵

$$\mathcal{D} = \{(x_i, y_i) \in \mathcal{X} \times \mathcal{Y} \mid i = 1, \dots, n\} = \{(X, Y)(\omega_i) \in \mathcal{X} \times \mathcal{Y} \mid i = 1, \dots, n\}$$

to estimate $\mathbb{P}X^{-1}$ by $\hat{\mathbb{P}}X^{-1} := \frac{1}{n} \sum_{i=1}^n \delta_{x_i}$, etc. When it comes to fairness, we extend the training data by the sensitive attribute:

$$\mathcal{D} = \{(x_i, s_i, y_i) \in \mathcal{X} \times \mathcal{S} \times \mathcal{Y} \mid i = 1, \dots, n\}.$$

Next to these distributions, often, a functional relation between the features X and the labels Y is the object of interest. This is done by learning the overall model \hat{Y} , composed of a learnable model (or model function) $f : \mathcal{X} \rightarrow \mathcal{Y}$, applied to the features $X : \Omega \rightarrow \mathcal{X}$. In such case, we consider a hypothesis space $\mathcal{H} = \{f : \mathcal{X} \rightarrow \mathcal{Y} \mid f \in \mathcal{H}\}$, i.e., a (sub-)space of functions mapping from the feature space \mathcal{X} to the target space \mathcal{Y} . Consecutively, we want to learn the relation between X and Y by finding the optimal function $f \in \mathcal{H}$, such that $\hat{Y} := f(X) \approx Y$ holds. In most cases, the hypothesis space is a set of functions $\mathcal{H} = \{f_\Theta : \mathcal{X} \rightarrow \mathcal{Y}, x \mapsto f_\Theta(x) \mid \Theta \in \mathbb{R}^{d_p}\}$ parameterized by a parameter $\Theta \in \mathbb{R}^{d_p}$. Finally, learning the functional relation between the features X and the labels Y by learning the optimal model $\hat{Y} \approx Y$ is done by comparing the results $\hat{y}_i = f_\Theta(x_i)$ of the model $\hat{Y} = f_\Theta(X)$ to the desirable results y_i for all $i = 1, \dots, n$ from the training data \mathcal{D} . The comparisons are done by using a suitable loss function which is applied to these magnitudes and optimized with respect to the parameter(s) that characterize(s) f_Θ .

Remark 2.4. Note that often, in ML-related literature, the introduction of Ω is omitted. Instead, random variables X on \mathcal{X} , Y on \mathcal{Y} , etc., are introduced. We introduce Ω to guarantee a well-defined usage of probabilities such as $\mathbb{P}(X = x)$ for some $x \in \mathcal{X}$, etc.

³Generalizations to multi-dimensional non-binary classification or regression tasks and multiple non-binary or continuous sensitive features are possible. While we will not go into detail regarding continuous labels or sensitive features, we will provide information on generalizability for the other cases.

⁴An arbitrary random variable $X : \Omega \rightarrow \mathcal{X}$ as given in subsection 2.1 induces a measure $\mathbb{P}X^{-1}$ on $\mathcal{F}_\mathcal{X}$ by $\mathbb{P}X^{-1}(A) := \mathbb{P}(X \in A)$ for all $A \in \mathcal{F}_\mathcal{X}$.

⁵Samples or statistical units can be both the elements $\omega \in \Omega$ from the population space as well as the realisations $(x, y) = (X, Y)(\omega) \in \mathcal{X} \times \mathcal{Y}$ of the random vector (X, Y) .

2.3 Generalized Notions of Group Fairness in Machine Learning

2.3.1 Motivation

Reflecting that there is no unique definition of fairness in real life, there is an enormous amount of different definitions of fairness in ML. While focusing on group fairness, even this category can be further grouped into three subcategories: Independence⁶, separation and sufficiency. In this context, group fairness can be characterized by some independence connected to the (binary) classification model $\hat{Y} : \Omega \rightarrow \mathcal{Y}$, the true label $Y : \Omega \rightarrow \mathcal{Y}$, and the sensitive feature $S : \Omega \rightarrow \mathcal{S}$. Barocas et al. [7] define these concepts as follows: *Independence* requires (mathematical) independence between the model’s prediction \hat{Y} and the sensitive feature S . *Separation* requires independence between the model’s prediction \hat{Y} and the sensitive feature S , conditioned on events based on the label Y . *Sufficiency* requires independence between the label Y and the sensitive feature S , conditioned on events based on the model’s prediction \hat{Y} . In this work, we will focus on the usually harder to achieve concepts of independence and separation.

However, there are also other definitions that fall under the broad umbrella of group fairness, but which can also be sorted in one of these subcategories (cf. [32]). They are usually defined for a one-dimensional binary classification task and for a single binary sensitive feature only, i.e., in settings where $\mathcal{Y} = \mathcal{S} = \{0, 1\}$ holds (cf. [26, 29, 32]). For example, one well-known fairness definition is called *disparate impact*. In most literature, it is assumed that $\{Y = 1\}$ is the class of interest and that $\{S = 0\}$ is the discriminated group, and that therefore, $\mathbb{P}(\hat{Y} = 1 \mid S = 0) < \mathbb{P}(\hat{Y} = 1 \mid S = 1)$ holds. In this case, the disparate impact score is defined as the proportion of the *passing rate* of the discriminated group from the privileged group

$$\text{DI} := \frac{\mathbb{P}(\hat{Y} = 1 \mid S = 0)}{\mathbb{P}(\hat{Y} = 1 \mid S = 1)} \quad (2.2)$$

and should satisfy $\text{DI} \geq 1 - \epsilon$ or $\text{DI} \geq \frac{p}{100}$ for some $\epsilon \in [0, 1]$ or $p \in [0, 100]$ (cf. [29]). The latter rule is also known as the $p\%$ -rule, and $p = 80$ (or $\epsilon = 0.2$) is a desirable choice (cf. [29, 43]). At the same time, the 80%-rule is also a popular legal term and the reason that the disparate impact score received its importance: It is “designed to mathematically represent the legal notion of *disparate impact*” (cf. [29]), which requires to avoid that “one group’s passing rate is less than 80% of the group with the highest rate” (cf. [10]).

The goal of the rest of this subsection is on the one hand to connect these different group fairness notions and on the other hand to introduce generalized notions for more general settings. More precisely, our contribution to this existing research is as follows:

- Starting from the definitions of Barocas et al. [7] for independence (subsection 2.3.2) and separation (subsection 2.3.3) each, we will make use of the particularly easy necessary and sufficient condition for the independence of two random variables (lemma 2.2) which we derived in subsection 2.1 to derive easily-to-test and generalized notions of group fairness in the context of subsection 2.2. In detail, these notions are applicable for more general settings, i.e., not only for a one-dimensional binary classification tasks and a single binary sensitive feature, which is the setting on which the majority of the literature focuses (cf. [16, 19, 21, 22, 26, 29, 32, 42, 43]).

⁶Mind the difference between the mathematical concept of independence and the fairness concept of independence. It is usually clear from the context what of both concepts is meant.

- Based on these notions, we will derive generalized empirical notions of the most common group fairness definitions...
- ... and prove that these coincide with the corresponding definitions in the setting of one-dimensional binary classification task and a single binary sensitive feature.

While the generalized empirical group fairness definitions will appear to be intrinsic compared to already existing definitions, our theoretical work in subsection 2.1 shows that these generalizations display not only a necessary, but a sufficient condition for the desired independence criterion they are based on. As a summary, an overview of already existing definitions and how we extend these is displayed in table 2.1 and 2.2 (subsection 2.3.4). For technical reasons, we assume that all of the following conditional probabilities exist.

2.3.2 Independence

An easily-to-test notion of group fairness

Definition 2.5 (Fairness according to the independence criterion (cf. [7])).

The classification model \hat{Y} is fair with respect to the sensitive feature S in the sense of the independence criterion if and only if (iff) \hat{Y} and S are mathematically independent with respect to \mathbb{P} .

Based on this definition of Barocas et al. [7], lemma 2.2 induces the following easily-to-test independence criterion in the context of fair ML:

Corollary 2.6 (Fairness according to the independence criterion).

The classification model \hat{Y} is fair with respect to the sensitive feature S in the sense of the independence criterion iff

$$\mathbb{P}(\hat{Y} = y \mid S = s_{k_1}) = \mathbb{P}(\hat{Y} = y \mid S = s_{k_2})$$

holds for all $y \in \mathcal{Y} = \{0, 1\} \subset \mathbb{N}_0$ and all $s_{k_1}, s_{k_2} \in \mathcal{S} = \{s_1, \dots, s_K\} \subset \mathbb{N}_0$.

Generalized empirical notions of group fairness In practise, exact equality according to corollary 2.6 is usually not achieved. This motivates to keep the difference of both sides of the equation(s) as small as possible, which translates to the following two *generalized* definitions of disparate impact and demographic parity.

More precisely, while for group fairness, the majority of the literature focuses on a binary classification model \hat{Y} and a single binary sensitive feature S (cf. [16, 19, 21, 22, 26, 29, 32, 42, 43]), in this work's definitions, we generalize the understanding of group fairness to a non-binary sensitive feature S , but which can also be used to model even multiple non-binary sensitive features (cf. remark 2.9).

While disparate impact is specifically designed for a binary classification task, i.e., for setting where $\mathcal{Y} = \{0, 1\}$ holds, and where the class $\{Y = 1\}$ is the preferred one (cf. remark 2.11), the demographic parity score additionally allows generalization to a one- or multidimensional non-binary classifier \hat{Y} by definition and based on the theoretical background in subsection 2.1⁷:

⁷If $\mathcal{Y} \neq \{0, 1\}$ is discrete, testing for the canonical one-elementary events $y \in \mathcal{Y}$ is still a necessary and sufficient condition for independence according to lemma 2.2.

Definition 2.7 (Disparate impact).

Let $\epsilon \in [0, 1]$. The disparate impact score

$$\text{DI} := \min_{s_{k_1}, s_{k_2} \in \mathcal{S}} \frac{\mathbb{P}(\hat{Y} = 1 \mid S = s_{k_1})}{\mathbb{P}(\hat{Y} = 1 \mid S = s_{k_2})}$$

measures the (un-)fairness of the classification model \hat{Y} with respect to the sensitive feature S in the sense of the independence criterion. For the model \hat{Y} , disparate impact is limited to ϵ iff $\text{DI} \geq 1 - \epsilon$ holds.

Definition 2.8 (Demographic parity).

Let $\epsilon \in [0, 1]$. The demographic parity score

$$\text{DP} := \max_{y \in \mathcal{Y}, s_{k_1}, s_{k_2} \in \mathcal{S}} \left| \mathbb{P}(\hat{Y} = y \mid S = s_{k_1}) - \mathbb{P}(\hat{Y} = y \mid S = s_{k_2}) \right|$$

measures the (un-)fairness of the classification model \hat{Y} with respect to the sensitive feature S in the sense of the independence criterion. For the model \hat{Y} , demographic parity holds with respect to ϵ iff $\text{DP} \leq \epsilon$ holds.

Remark 2.9. In our previous work [35], we consider $K = |\mathcal{S}|$ different binary random variables $S_1, \dots, S_K : \Omega \rightarrow \{0, 1\}$ when defining disparate impact. Encoding the single non-binary random variable S from this work to K such binary random variables S_1, \dots, S_K for all $k = 1, \dots, K$ yields the same definition of disparate impact as given in [35]. We change the notation in this work because it is more intuitive compared to common fairness definitions (e.g., cf. eq. (2.2) and proof of lemma 2.10) and easily shows how these fairness definitions can be extended even to *multiple* non-binary sensitive features: In this case, the random *vector* $S = (S_1, \dots, S_{d_s}) : \Omega \rightarrow \mathcal{S}$ with $\mathcal{S} = \mathcal{S}_1 \times \dots \times \mathcal{S}_{d_s} \subset \mathbb{N}_0^{d_s}$ and $d_s > 1$ encodes all d_s possibly non-binary single sensitive features S_l for $l = 1, \dots, d_s$.

Accordance of empirical notions of group fairness in the binary case In case of a binary classification task and a single binary sensitive feature, our definitions coincide with the according definitions known from the before-mentioned literature:

Lemma 2.10. *If $\mathcal{Y} = \mathcal{S} = \{0, 1\}$ holds, the disparate impact score DI and the demographic parity score DP according to definition 2.7 and 2.8, respectively, coincide with the corresponding definitions known from the literature.*

Proof. If $\mathcal{Y} = \mathcal{S} = \{0, 1\}$ holds, the fact that $\{\hat{Y} = 0\} \cup \{\hat{Y} = 1\} = \Omega$ holds implies that the probability measure $\mathbb{P}(\hat{Y} \in \cdot \mid S = s)$ is uniquely determined by the probability $\mathbb{P}(\hat{Y} = 1 \mid S = s)$ for all $s \in \mathcal{S}$. Therefore, the independence criterion (corollary 2.6) becomes

$$\mathbb{P}(\hat{Y} = 1 \mid S = 0) = \mathbb{P}(\hat{Y} = 1 \mid S = 1).$$

By the same fact,

$$\begin{aligned} & |\mathbb{P}(\hat{Y} = 0 \mid S = 0) - \mathbb{P}(\hat{Y} = 0 \mid S = 1)| \\ &= |1 - \mathbb{P}(\hat{Y} = 1 \mid S = 0) - (1 - \mathbb{P}(\hat{Y} = 1 \mid S = 1))| \\ &= |\mathbb{P}(\hat{Y} = 1 \mid S = 0) - \mathbb{P}(\hat{Y} = 1 \mid S = 1)| \end{aligned}$$

holds. Therefore, the demographic parity score (definition 2.8) becomes

$$DP = \left| \mathbb{P}(\hat{Y} = 1 \mid S = 0) - \mathbb{P}(\hat{Y} = 1 \mid S = 1) \right|.$$

Moreover, the disparate impact score (definition 2.7) becomes

$$DI = \min \left\{ \frac{\mathbb{P}(\hat{Y} = 1 \mid S = 0)}{\mathbb{P}(\hat{Y} = 1 \mid S = 1)}, \frac{\mathbb{P}(\hat{Y} = 1 \mid S = 1)}{\mathbb{P}(\hat{Y} = 1 \mid S = 0)} \right\}.$$

In most literature, where $\{S = 0\}$ is assumed to be the discriminated group, and therefore, where $\mathbb{P}(\hat{Y} = 1 \mid S = 0) < \mathbb{P}(\hat{Y} = 1 \mid S = 1)$ holds, this simplifies to

$$DI = \frac{\mathbb{P}(\hat{Y} = 1 \mid S = 0)}{\mathbb{P}(\hat{Y} = 1 \mid S = 1)}$$

(cf. eq. (2.2)). These are the definitions of the disparate impact and the demographic parity score usually found in the literature (cf. [26, 29, 32, 43]). \square

As already briefly touched on in the motivation subsection 2.3.1, the disparate impact criterion assures that the relative amount of positive predictions within the discriminated group $\{S = 0\}$ – or in our generalized case of non-binary sensitive features, within the most discriminated group – deviates at most $(100-p)\% = 100\epsilon\%$ from the relative amount of positive predictions within the privileged group $\{S = 1\}$ – or in our generalized case, within the most privileged group (cf. definition 2.7). For short and in either way: It aims at obtaining similar or equal *success* or *passing rates* among groups.

Similarly, in a binary classification task, the demographic parity criterion assures that the relative amount of positive predictions deviates at most $100\epsilon\%$ among groups (cf. proof of lemma 2.10 or table 2.2). In contrast, in a non-binary classification task, the demographic parity criterion assures that the relative amount of *any* predictions deviates at most $100\epsilon\%$ among groups (cf. definition 2.8).

By that, while both criteria assure similar or equal passing rates among groups in the setting of a binary classification task, they assure different things in the setting of a non-binary classification task due to the consideration of all labels in the demographic parity criterion (cf. table 2.2).

Remark 2.11 (Generalizability of the disparate impact score). Similar to the demographic parity score DP (definition 2.8), one could ask whether it makes sense to generalize the disparate impact score DI to arbitrary discrete target spaces \mathcal{Y} by

$$DI^* := \min_{y \in \mathcal{Y}, s_{k_1}, s_{k_2} \in \mathcal{S}} \frac{\mathbb{P}(\hat{Y} = y \mid S = s_{k_1})}{\mathbb{P}(\hat{Y} = y \mid S = s_{k_2})}. \quad (2.3)$$

However, this generalized definition would *not* coincide with the common one from equation (2.2) in the setting of lemma 2.10: For example, consider the case

$$\begin{aligned} \mathbb{P}(\hat{Y} = 1 \mid S = 0) &= 0.8, & \mathbb{P}(\hat{Y} = 1 \mid S = 1) &= 0.9, \\ \mathbb{P}(\hat{Y} = 0 \mid S = 0) &= 0.2, & \mathbb{P}(\hat{Y} = 0 \mid S = 1) &= 0.1. \end{aligned}$$

In this case, the disparate impact score according to definition 2.7 is equal to $DI = \min\{\frac{0.8}{0.9}, \frac{0.9}{0.8}\} = \frac{0.8}{0.9} \approx 0.88$, which usually is a score considered to be fair. In contrast, the disparate impact score according to equation (2.3) is equal to $DI^* = \min\{DI, \frac{0.2}{0.1}, \frac{0.1}{0.2}\} = \frac{0.1}{0.2} = 0.5$, which usually is a score considered to be unfair. The reason is that the idea of disparate impact relies on the fact that the class $\{Y = 1\}$ is the desired one and only the relative amount of positive predictions among groups is of interest (cf. [29]). Therefore, it does only make sense to define disparate impact score as we do in definition 2.7.

2.3.3 Separation

An easily-to-test notion of group fairness Depending on the case of application, one disadvantage of fairness notions that belong to the fairness concept *independence* could be the missing dependence on the true label Y . In such case, even if the model \hat{Y} was perfect, i.e., if $\hat{Y} = Y$ held, it would be denoted as unfair if the relative amount of positive training labels differed significantly among groups (cf. [19]).

Solution to that yields the fairness concept *separation*, which in contrast to the fairness concept independence requires (mathematical) independence between the model's prediction \hat{Y} and the sensitive feature S , conditioned on Y :

Definition 2.12 (Fairness according to the separation criterion (cf. [7])).

The classification model \hat{Y} is fair with respect to the sensitive feature S in the sense of the separation criterion iff \hat{Y} and S are mathematically independent with respect to $\mathbb{P}(\cdot \mid Y = y)$ for all $y \in \mathcal{Y}$.

Using the modified probability measure $\mathbb{P}(\cdot \mid Y = y)$ for $y \in \mathcal{Y}$, lemma 2.2 again induces the following easily-to-test separation criterion in the context of fair ML:

Corollary 2.13 (Fairness according to the separation criterion).

The classification model \hat{Y} is fair with respect to the sensitive feature S in the sense of the separation criterion iff

$$\mathbb{P}(\hat{Y} = \hat{y} \mid S = s_{k_1}, Y = y) = \mathbb{P}(\hat{Y} = \hat{y} \mid S = s_{k_2}, Y = y)$$

holds for all $y, \hat{y} \in \mathcal{Y} = \{0, 1\} \subset \mathbb{N}_0$ and all $s_{k_1}, s_{k_2} \in \mathcal{S} = \{s_1, \dots, s_K\} \subset \mathbb{N}_0$.

Generalized empirical notions of group fairness Again, in practise, exact equality according to corollary 2.13 is usually not achieved. Therefore, again, keeping the difference of both sides of the equation as small as possible motivates the following *generalized* definitions, where similar to the previous subsection 2.3.2, the second one is specifically designed for a binary classification task, i.e., for settings where $\mathcal{Y} = \{0, 1\}$ holds, and where the class $\{Y = 1\}$ is the preferred one.

Definition 2.14 (Equalized odds).

Let $\epsilon \in [0, 1]$. The equalized odds scores

$$\text{EOs}(y) := \max_{s_{k_1}, s_{k_2} \in \mathcal{S}} \left| \mathbb{P}(\hat{Y} = y \mid S = s_{k_1}, Y = y) - \mathbb{P}(\hat{Y} = y \mid S = s_{k_2}, Y = y) \right|$$

measure the (un-)fairness of the classification model \hat{Y} with respect to the sensitive feature S in the sense of the separation criterion for all $y \in \mathcal{Y}$. For the model \hat{Y} , equalized odds holds with respect to ϵ iff $\text{EOs}(y) \leq \epsilon$ holds for all $y \in \mathcal{Y}$.⁸

Definition 2.15 (Equal opportunity).

Let $\epsilon \in [0, 1]$. The equal opportunity score

$$\text{EO} := \max_{s_{k_1}, s_{k_2} \in \mathcal{S}} \left| \mathbb{P}(\hat{Y} = 1 \mid S = s_{k_1}, Y = 1) - \mathbb{P}(\hat{Y} = 1 \mid S = s_{k_2}, Y = 1) \right|$$

measures the (un-)fairness of the classification model \hat{Y} with respect to the sensitive feature S in the sense of the separation criterion. For the model \hat{Y} , equal opportunity holds with respect to ϵ iff $\text{EO} \leq \epsilon$ holds.

⁸Instead of the requirement “for all $y \in \mathcal{Y}$ ”, we could also take the maximum over $y \in \mathcal{Y}$ and $s_{k_1}, s_{k_2} \in \mathcal{S}$ and consider a single equalized odds score as done in definition 2.8.

Similar arguments as compared to subsection 2.3.2 also show how these definition(s) allow a generalized understanding of group fairness for non-binary and even multiple non-binary sensitive features S , and for a one- or multi-dimensional non-binary classifier \hat{Y} .

Accordance of empirical notions of group fairness in the binary case In case of a binary classification task and a single binary sensitive feature, our definitions coincide with the according definitions known from other literature:

Lemma 2.16. *If $\mathcal{Y} = \mathcal{S} = \{0, 1\}$ holds, the equalized odds scores $\text{EOs}(y)$ for $y \in \mathcal{Y}$ and the equal opportunity score EO according to definition 2.14 and 2.15, respectively, coincide with the corresponding definitions known from the literature.*

Proof. If $\mathcal{Y} = \mathcal{S} = \{0, 1\}$ holds, similar to the proof of lemma 2.10, the separation criterion (definition 2.13) becomes

$$\mathbb{P}(\hat{Y} = 1 \mid S = 0, Y = y) = \mathbb{P}(\hat{Y} = 1 \mid S = 1, Y = y)$$

for $y = 0, 1$, and

$$\begin{aligned} & |\mathbb{P}(\hat{Y} = 0 \mid S = 0, Y = 0) - \mathbb{P}(\hat{Y} = 0 \mid S = 1, Y = 0)| \\ &= |1 - \mathbb{P}(\hat{Y} = 1 \mid S = 0, Y = 0) - (1 - \mathbb{P}(\hat{Y} = 1 \mid S = 1, Y = 0))| \\ &= |\mathbb{P}(\hat{Y} = 1 \mid S = 0, Y = 0) - \mathbb{P}(\hat{Y} = 1 \mid S = 1, Y = 0)| \end{aligned}$$

holds.⁹ Therefore, the equalized odds scores (definition 2.14) become

$$\begin{aligned} \text{EOs}(1) &= |\mathbb{P}(\hat{Y} = 1 \mid S = 0, Y = 1) - \mathbb{P}(\hat{Y} = 1 \mid S = 1, Y = 1)| \quad \text{and} \\ \text{EOs}(0) &= |\mathbb{P}(\hat{Y} = 1 \mid S = 0, Y = 0) - \mathbb{P}(\hat{Y} = 1 \mid S = 1, Y = 0)| \end{aligned}$$

(comparison of TPRs and FPRs among groups) and the equal opportunity score (definition 2.15) becomes

$$\text{EO} = |\mathbb{P}(\hat{Y} = 1 \mid S = 0, Y = 1) - \mathbb{P}(\hat{Y} = 1 \mid S = 1, Y = 1)|$$

(comparison of TPRs among groups). These are the definitions of the equalized odds and the equal opportunity score(s) usually found in the literature (cf. [26, 29, 32, 42]). \square

While equalized odds ensures that the true positive rates (TPRs) *and* true negative rates (TNRs) (or equivalently, false positive rates (FPRs)) among groups differ at most 100ε% in a binary classification task, equalized odds only concentrates on TPRs among groups. In contrast, in a non-binary classification task where the TPRs and FPRs are not well-defined, equalized odds refers to similar or equal *correct classification rates* per label among groups (cf. definition 2.14) and displays a natural generalization of equal opportunity in this setting (cf. definition 2.15).

Remark 2.17. Nevertheless, we will not make use of equalized odds in this work, as the TNRs and FPRs given by $\mathbb{P}(\hat{Y} = 0 \mid S = s, Y = 0)$ and $\mathbb{P}(\hat{Y} = 1 \mid S = s, Y = 0)$ for any $s \in \mathcal{S}$, respectively, do not exist in our domain of application although being in the setting of a binary classification task, as we will see in subsection 3.2.

⁹More precisely, comparing the TNRs is equivalent to comparing the FPRs.

2.3.4 Summary of Generalized Notions of Group Fairness

To conclude, in this section, we derived generalized exact and empirical notions of group fairness based on the mathematical concept of independence and suitable for a single, but also multiple non-binary sensitive feature(s). All exact and some empirical notions are suitable for not only one-, but also multi-dimensional non-binary classification models. We additionally showed that the notions coincide with common group fairness definitions in the case of a binary classification task and a single binary sensitive feature. A summary of such already existing definitions and our contributions are summarized in table 2.1 and 2.2.

	Definition according to Barocas et al.	Derived necessary and sufficient condition
Independence	$\hat{Y} \perp\!\!\!\perp S$	$\mathbb{P}(\hat{Y} = y \mid S = s_{k_1})$ $=$ $\mathbb{P}(\hat{Y} = y \mid S = s_{k_2})$ $\forall y \in \mathcal{Y}, s_{k_1}, s_{k_2} \in \mathcal{S}$
Separation	$\hat{Y} \perp\!\!\!\perp S \mid Y$	$\mathbb{P}(\hat{Y} = \hat{y} \mid S = s_{k_1}, Y = y)$ $=$ $\mathbb{P}(\hat{Y} = \hat{y} \mid S = s_{k_2}, Y = y)$ $\forall y, \hat{y} \in \mathcal{Y}, s_{k_1}, s_{k_2} \in \mathcal{S}$

Table 2.1: Overview of exact fairness definitions. Overview of our exact derived necessary, sufficient and easily-to-test fairness conditions (corollary 2.6 and 2.13) based on the corresponding definitions of Barocas et al. [7] (definition 2.5 and 2.12).

	Derived generalized empirical definitions (multi cases)	Existing empirical definitions (binary cases)
Ind.: DI	$\mathcal{Y} = \{0, 1\}, \mathcal{S} \text{ arbitrary:}$ $\min_{s_{k_1}, s_{k_2} \in \mathcal{S}} \frac{\mathbb{P}(\hat{Y}=1 \mid S=s_{k_1})}{\mathbb{P}(\hat{Y}=1 \mid S=s_{k_2})}$	$\mathcal{Y} = \mathcal{S} = \{0, 1\}:$ $\frac{\mathbb{P}(\hat{Y}=1 \mid S=0)}{\mathbb{P}(\hat{Y}=1 \mid S=1)}$
Ind.: DP	$\mathcal{Y}, \mathcal{S} \text{ arbitrary:}$ $\max_{y \in \mathcal{Y}, s_{k_1}, s_{k_2} \in \mathcal{S}} \left \mathbb{P}(\hat{Y} = y \mid S = s_{k_1}) - \mathbb{P}(\hat{Y} = y \mid S = s_{k_2}) \right $	$\mathcal{Y} = \mathcal{S} = \{0, 1\}:$ $\left \mathbb{P}(\hat{Y} = 1 \mid S = 0) - \mathbb{P}(\hat{Y} = 1 \mid S = 1) \right $
Sep.: EO	$\mathcal{Y} = \{0, 1\}, \mathcal{S} \text{ arbitrary:}$ $\max_{s_{k_1}, s_{k_2} \in \mathcal{S}} \left \mathbb{P}(\hat{Y} = 1 \mid S = s_{k_1}, Y = 1) - \mathbb{P}(\hat{Y} = 1 \mid S = s_{k_2}, Y = 1) \right $	$\mathcal{Y} = \mathcal{S} = \{0, 1\}:$ $\left \mathbb{P}(\hat{Y} = 1 \mid S = 0, Y = 1) - \mathbb{P}(\hat{Y} = 1 \mid S = 1, Y = 1) \right $
Sep.: EOs	$\mathcal{Y}, \mathcal{S} \text{ arbitrary, } \forall y \in \mathcal{Y}:$ $\max_{s_{k_1}, s_{k_2} \in \mathcal{S}} \left \mathbb{P}(\hat{Y} = y \mid S = s_{k_1}, Y = y) - \mathbb{P}(\hat{Y} = y \mid S = s_{k_2}, Y = y) \right $	$\mathcal{Y} = \mathcal{S} = \{0, 1\}:$ $\left \mathbb{P}(\hat{Y} = 1 \mid S = 0, Y = 1) - \mathbb{P}(\hat{Y} = 1 \mid S = 1, Y = 1) \right $ $\left \mathbb{P}(\hat{Y} = 1 \mid S = 0, Y = 0) - \mathbb{P}(\hat{Y} = 1 \mid S = 1, Y = 0) \right $

Table 2.2: Overview of empirical fairness definitions. Comparison of our generalized empirical fairness definitions (definition 2.7, 2.8, 2.14 and 2.15) and the corresponding existing definitions (e.g., cf. [29]).

Remark 2.18 (Computation of group fairness scores in practise). In practise, the true distributions $\mathbb{P}(S, \hat{Y})^{-1}$ and $\mathbb{P}(S, Y, \hat{Y})^{-1}$, on which the probabilities displayed in table 2.1

and 2.2 are based on, are unknown. Therefore, as elaborated in subsection 2.2, the fairness scores are computed using the empirical approximations $\hat{\mathbb{P}}(S, \hat{Y})^{-1} = \frac{1}{n} \sum_{i=1}^n \delta_{(s_i, \hat{y}_i)}$ and $\hat{\mathbb{P}}(S, Y, \hat{Y})^{-1} = \frac{1}{n} \sum_{i=1}^n \delta_{(s_i, y_i, \hat{y}_i)}$ based on the training data \mathcal{D} , respectively, yielding the required approximated probabilities

$$\begin{aligned} \hat{\mathbb{P}}(\hat{Y} = y \mid S = s) &= \frac{\sum_{i=1}^n \mathbb{1}_{\{\hat{y}_i = y, s_i = s\}}}{\sum_{i=1}^n \mathbb{1}_{\{s_i = s\}}} & \text{and} \\ \mathbb{P}(\hat{Y} = \hat{y} \mid S = s, Y = y) &= \frac{\sum_{i=1}^n \mathbb{1}_{\{\hat{y}_i = \hat{y}, s_i = s, y_i = y\}}}{\sum_{i=1}^n \mathbb{1}_{\{s_i = s, y_i = y\}}} & \text{for all } y, \hat{y} \in \mathcal{Y}, s \in \mathcal{S}. \end{aligned}$$

3 Leakage Detection in Water Distribution Systems

In view of the AI Act, by being part of the critical infrastructure, WDSs belong to *high-risk systems* (cf. [38]). In this context, “[m]uch attention has been paid to the potential for AI systems to facilitate indirect discrimination, [which is] in principle illegal under EU law” ([38]). One requirement of such systems is therefore to check the system for bias and to document the system’s performance for different demographic groups (cf. [36]). While this could suggest the use of group fairness definitions implicitly, the guidelines for trustworthy AI explicitly name fairness as one of the seven essential requirements for such systems (cf. [15]).

A key challenge in the domain of WDSs where AI, or more precisely, ML, is used, is to detect leakages (cf. [5, 18, 25, 31, 34, 39]). The main components of a WDS relevant for this work are *nodes* and *pipes*, through which water can be supplied to end users such as private households, hospitals or schools located at the nodes of the network, but which are also vulnerable to leakages. To detect such is therefore crucial to guarantee consistent water supply, but can also affect other important tasks such as short-term decision making in and long-term planning of WDSs.

Therefore, as requested by the AI Act and the guidelines for trustworthy AI, in this section, we present a common ML-based pipeline and concrete instantiations of how to detect leakages in WDSs (subsection 3.1). Consecutively, we investigate what fairness can mean (subsection 3.2) and whether it is satisfied (subsection 3.3 and 3.4) in this context and according to common group fairness notions as introduced in subsection 2.3.

3.1 Methodology

In the task of leakage detection, the domain Ω (cf. subsection 2.2) corresponds to possible states of a WDS, determined by time-dependent demands of the end users located at the $D \in \mathbb{N}$ nodes in the system and which may be affected by leakages. We assume that among those nodes, $d \in \mathbb{N}$ nodes are provided with sensors (usually, $D \gg d$), which deliver pressure measurements $p(t) \in \mathbb{R}^d$ at different times $t \in \mathbb{R}$ and which can be used for the task at hand. As the sensors usually measure pressure values within fixed time intervals $\delta \in \mathbb{R}_+$, we introduce the notation $t_i := t_0 + i\delta$, where t_0 is some fixed reference point with respect to time.

There are several methodologies that make use of such pressure measurements to approach the problem of leakage detection using ML. Using the notation from subsection 2.2, the goal is to train a binary classifier $\hat{Y} = f(X) : \Omega \rightarrow \mathcal{X} \rightarrow \mathcal{Y}$ with $\mathcal{Y} = \{0, 1\}$ that predicts the true state $Y : \Omega \rightarrow \mathcal{Y}$ of the WDS with respect to the question whether a leakage is

active ($\{Y = 1\}$) or not ($\{Y = 0\}$). Hereby, the feature space \mathcal{X} depends on the specific method, but is related to the before-mentioned pressure measurements.

One standard approach comes in three steps (cf. [20]): First of all, so called *virtual sensors* are trained, i.e., regression models that are able to predict the pressure at some time $t_i \in \mathbb{R}$ and at a node $j \in \{1, \dots, d\}$ (or even $j \in \{1, \dots, D\}$), based on the pressure measurements observed at that (or earlier) time(s) and at (a choice of) the sensor nodes $j \in \{1, \dots, d\}$. Subsequently, these virtual sensors are used to compute *pressure residuals* of measured and predicted pressure. Finally, these pressure residuals are fed into a *leakage detector* \hat{Y} that is able to predict whether a leakage is present in the WDS at the time of the used residual (cf. [20]). An overview of this pipeline is displayed in figure 3.1.

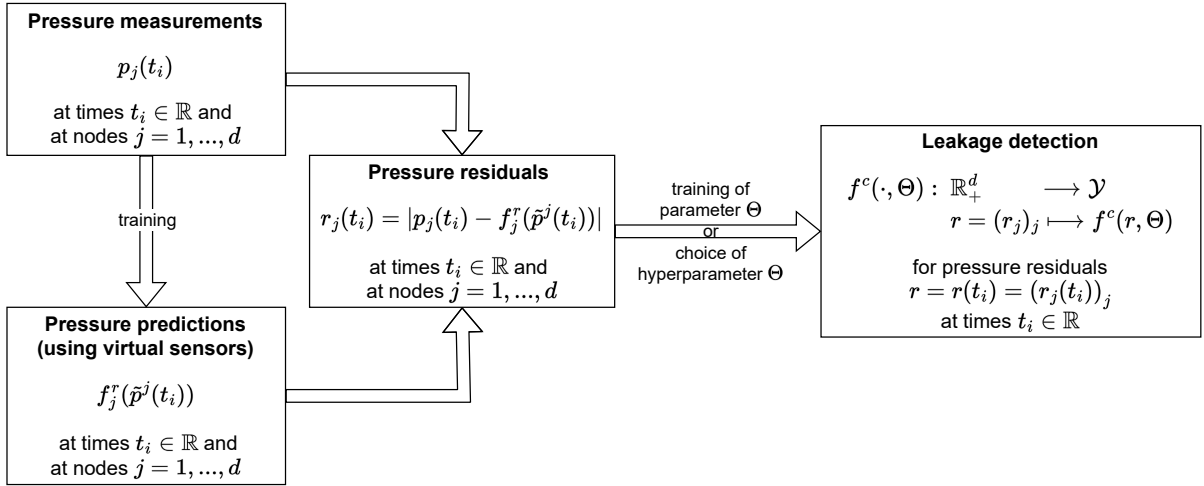


Figure 3.1: Standard leakage detection pipeline.

The approach can differ depending on the concrete instantiation of virtual sensors and the leakage detector. In this subsection, we first formalize the idea of the general leakage detection pipeline described above in more detail (subsection 3.1.1). Consecutively, we present two concrete instantiations of such (subsection 3.1.2), which we will investigate on with respect to the question of fairness in the rest of this section.

3.1.1 Leakage Detection Pipeline

Virtual sensors Based on vector inputs $\tilde{p}^j(t_i) \in \mathbb{R}^{d_r}$ that are based on the pressure measurements $p(t_\ell) = (p_j(t_\ell))_{j=1, \dots, d} \in \mathbb{R}^d$ observed at (multiple) times $t_\ell \in \mathbb{R}$ and at the sensor nodes $j = 1, \dots, d$ in the WDS, so called *virtual sensors*, i.e., regression models

$$f_j^r : \mathbb{R}^{d_r} \longrightarrow \mathbb{R}$$

$$\tilde{p}^j(t_i) \longmapsto f_j^r(\tilde{p}^j(t_i))$$

that predict the pressure at times $t_i \in \mathbb{R}$ and at the sensor node j are trained for each sensor node $j = 1, \dots, d$. Hereby, the dimension $d_r \in \mathbb{N}$ and the inputs $\tilde{p}^j(t_i) \in \mathbb{R}^{d_r}$ depend on the specific model architecture used (cf. subsection 3.1.2 and [5, 6, 20]).

Pressure residuals Independent of what specific instantiations of virtual sensors f_j^r for $j = 1, \dots, d$ are used, standard leakage detection methods rely on the *pressure residuals*

$$r_j(t_i) := |p_j(t_i) - f_j^r(\tilde{p}^j(t_i))| \in \mathbb{R}_+$$

we obtain from the pressure measurements $p_j(t_i) \in \mathbb{R}$ and the pressure predictions $f_j^r(\tilde{p}^j(t_i)) \in \mathbb{R}$ at (possibly unseen) times $t_i \in \mathbb{R}$ and at the sensor node j for all $j = 1, \dots, d$ (cf. [5, 20]).

Leakage detection Based on pressure residuals $r(t_i) = (r_j(t_i))_{j=1,\dots,d} \in \mathcal{X} := \mathbb{R}^{d_c} = \mathbb{R}^d$ (i.e., $d_c = d$) at times $t_i \in \mathbb{R}$ and at the sensor nodes $j = 1, \dots, d$ in the WDS, a classification model \hat{Y} – or more precisely, the learnable model f^c which is applied to the feature pressure residuals $X : \Omega \rightarrow \mathcal{X}$ (cf. subsection 2.2) –

$$f^c := f_\Theta^c := f^c(\cdot, \Theta) : \mathcal{X} \longrightarrow \mathcal{Y}$$

that predicts whether ($\{\hat{Y} = f^c(X) = 1\}$) or whether not ($\{\hat{Y} = f^c(X) = 0\}$) a leakage is present in the WDS is defined or trained. Hereby, $\Theta \in \mathbb{R}^{d_p}$ indicates a choosable or trainable (hyper-)parameter and the hypothesis space \mathcal{H} depends on the specific model architecture used (cf. subsection 3.1.2, section 4 and [5, 20]).

3.1.2 Leakage Detection Instantiations

The previous subsection gives a general pipeline on how to detect leakages in a WDS based on pressure measurements, pressure predictions based on virtual sensors, resulting pressure residuals and finally, the leakage detection itself (cf. figure 3.1). In this subsection, we present specific instantiations of this approach.

Linear virtual sensors The first approach is based on the work of Artelt et al. [5]: In this case, each virtual sensor $f_j^r : \mathbb{R}^{d_r} \rightarrow \mathbb{R}$ at each sensor node $j \in \{1, \dots, d\}$ corresponds to a linear regression model. The inputs $\tilde{p}^j(t_i) \in \mathbb{R}^{d_r}$ at times $t_i \in \mathbb{R}$ consist of rolling means

$$\tilde{p}^j(t_i) := \bar{p}_{\neq j}(t_i) := \frac{1}{T_r + 1} \sum_{\iota=0}^{T_r} p_{\neq j}(t_i - \iota\delta) \in \mathbb{R}^{d-1}$$

at all sensor nodes except the node j and with a to be chosen time window $T_r + 1 \in \mathbb{N}$. By that, each regression model's input dimension equals $d_r := d - 1$. Based on that, the d virtual sensors f_j^r for each sensor node $j \in \{1, \dots, d\}$ are trained on leakage free training data $\mathcal{D}_j^r = \{(\bar{p}_{\neq j}(t_i), p_j(t_i)) \in \mathbb{R}^{d_r} \times \mathbb{R} \mid i = 0, \dots, n_r\}$. More precisely, $y(t_i) = 0 \in \mathcal{Y}$ holds for all realisations $i = 0, \dots, n_r$ of the label Y .

GCN virtual sensors In contrast, the second approach is based on the work of Ashraf et al. [6]: In this case, each virtual sensor $f_j^r : \mathbb{R}^{d_r} \rightarrow \mathbb{R}$ at each sensor node $j \in \{1, \dots, d\}$ is obtained by training a single GCN model.

The GCN model is trained on leakage free training data

$$\mathcal{D}^r = \left\{ \left((p_j(t_i))_{j=1,\dots,d}, (p_j(t_i))_{j=1,\dots,D} \right) \in \mathbb{R}^d \times \mathbb{R}^D \mid i = 0, \dots, n_r \right\}.$$

More precisely, the GCN model inputs the sparse pressure measurements at the sensor nodes $j = 1, \dots, d$ and outputs the pressure predictions at *each* node $j = 1, \dots, D$ of the WDS. However, for this work, the pressure predictions at the sensor nodes $j = 1, \dots, d$ are enough: The d virtual sensors f_j^r at each sensor node $j \in \{1, \dots, d\}$ can be considered as the entry-wise output of the overall GCN model $f^r := (f_j^r)_{j=1,\dots,D}$.

By that, the inputs $\tilde{p}^j(t_i) \in \mathbb{R}^{d_r}$ at times $t_i \in \mathbb{R}$ are given by the node-independent pressure measurements $\tilde{p}(t_i) = \tilde{p}^j(t_i) := (p_j(t_i))_{j=1,\dots,d} \in \mathbb{R}^d$ themselves for *all* sensor nodes $j = 1, \dots, d$, and each regression model's input dimension equals $d_r := d$.

Ensemble leakage detection: The H-method Independent on the choice of virtual sensor, based on the pressure residuals $r(t_i) = (r_j(t_i))_{j=1,\dots,d} \in \mathcal{X} = \mathbb{R}_+^{d_c} = \mathbb{R}_+^d$ at times $t_i \in \mathbb{R}$ we obtain from these, a simple leakage detection method performing good on standard benchmarks is the *threshold-based ensemble classification* introduced by Artelt et al. [5]: Without any further training, we can choose a node-dependent hyperparameter $\theta_j \in \mathbb{R}_+$ to define a (local) classifier $f_j^c : \mathbb{R}_+ \rightarrow \mathcal{Y}$ for each sensor node $j \in \{1, \dots, d\}$ by

$$f_j^c(r_j(t_i)) = f_j^c(r_j(t_i), \theta_j) := \mathbb{1}_{\{r_j(t_i) > \theta_j\}}.$$

We then obtain an ensemble classifier $f^c : \mathcal{X} \rightarrow \mathcal{Y}$ with feature space $\mathcal{X} = \mathbb{R}_+^{d_c} = \mathbb{R}_+^d$ and hyperparameter $\Theta := (\theta_j)_{j=1,\dots,d_c} \in \mathcal{X}$ (i.e., $d_p = d_c = d$) that predicts whether there is a leakage present in the WDS at time $t_i \in \mathbb{R}$ or not, defined by

$$f^c(r(t_i)) = f^c(r(t_i), \Theta) := \mathbb{1}_{\left\{\sum_{j=1}^{d_c} f_j^c(r_j(t_i)) \geq 1\right\}}. \quad (3.4)$$

Simply put into words, a node-dependent classifier f_j^c detects a leakage when the node-dependent pressure-residual $r_j(t_i) \in \mathbb{R}_+$ at time $t_i \in \mathbb{R}$ exceeds the node-dependent threshold $\theta_j \in \mathbb{R}_+$ and the ensemble classifier f^c detects a leakage when one of the node-dependent classifiers f_j^c for any $j \in \{1, \dots, d\}$ does.

We call this overall instantiation of the standard leakage detection pipeline (cf. figure 3.1) independently on the instantiation of the virtual sensors and characterized by choosing the **Hyperparameter** $\Theta \in \mathcal{X}$ the **H-method**. Note that the H-method does not need further training once it has access to feature pressure residuals $\mathcal{X} = \mathbb{R}_+^d$. How to introduce a trainable structure to this last component of the pipeline will be part of section 4.

3.2 Fairness in Leakage Detection

After having introduced a pipeline to define a leakage detection model $\hat{Y} = f^c(X)$ and possible concrete instantiations of such in the previous subsection, the question arises how the leakage detection is related to fairness in the sense of subsection 2.3. One key contribution of this work is to answer this question, i.e., to introduce the notion of fairness in the application domain of WDSs by defining suitable sensitive features in the context of leakage detection or other ML-based services in WDSs.

Sensitive features in ML-based services in WDS Knowing that each node of the WDS corresponds to a group of consumers, a natural question is whether these local groups benefit from the WDS and its related services, such as the leakage detection, in equal degree. To ensure that the methods that will be presented in subsection 4.1 scale to larger WDSs, we do not consider single nodes but $K \in \mathbb{N}$ groups of nodes in the WDS as protected groups in terms of fairness. Then, given that a leakage is active in the WDS, i.e., that $Y = 1$ holds, we define the sensitive feature $S \in \mathcal{S} := \{s_1, \dots, s_K\} := \{1, \dots, K\}$ to give answer to the question where, i.e., in which protected group $k \in \{1, \dots, K\}$, this leakage is active.¹⁰ In terms of equal service one would expect an equally good detection of leakages independent on the leakage location, i.e., the protected group. This understanding of sensitive features, protected groups and consecutively, fairness in WDS, can of course be adapted to other ML-based services in WDS, for example, to contamination detection.

¹⁰Note that in contrast to common settings, where the random variables Y, \hat{Y} and S share the same domain, with respect to this fairness question, for the sensitive feature S , we change the domain from the population space Ω of all possible states to those, in which a leakage is present, i.e., to $\{Y = 1\} \subset \Omega$.

Fairness notions in ML-based services in WDS In this work, we will focus on the evaluation of fairness by choosing one fairness notion each from the fairness concepts independence and separation (cf. subsection 2.3.2 and 2.3.3): Disparate impact for independence (definition 2.7) due to its importance also in the legal context (cf. subsection 2.3.1), and equal opportunity for separation (definition 2.15). Regarding the latter concept, considering that our sensitive feature S is defined on the event $\{Y = 1\}$, this shows why using equalized odds is not possible in this setting, as already mentioned in remark 2.17 and as shown in the proof of the next lemma 3.1.

Fairness properties in ML-based services in WDS Given *this* definition of a non-binary sensitive feature S in the WDS, we obtain the following important results with regard to the notions of fairness chosen.

Lemma 3.1 (Equivalence of disparate impact and equal opportunity in WDSs). *Let $S : \{Y = 1\} \rightarrow \{1, \dots, K\}$ be the sensitive feature describing where a leakage in one of the protected groups $k \in \{1, \dots, K\}$ of the WDS is active. Moreover, let $\epsilon, \tilde{\epsilon} \in [0, 1]$ and define $\max_k := \max_{k \in \{1, \dots, K\}} \mathbb{P}(\hat{Y} = 1 \mid S = k)$.*

1. *If disparate impact is limited to ϵ , equal opportunity holds with respect to $\tilde{\epsilon} = \epsilon \max_k$.*
2. *If equal opportunity holds with respect to $\tilde{\epsilon}$, disparate impact is limited to $\epsilon = \tilde{\epsilon}(\max_k)^{-1}$.*

Proof. First of all, note that for any $\omega \in \Omega$ for which there exists a $k \in \{1, \dots, K\}$ such that $S(\omega) = s_k = k$ holds, $Y(\omega) = 1$ must hold by definition of the sensitive feature S (this is why it only makes sense to define S on $\{Y = 1\} \subset \Omega$). Therefore, $\{S = k, Y = 0\}$ is empty for all $k = 1, \dots, K$. Subsequently, we obtain

$$\begin{aligned} \{S = k, Y = 1\} &= \{S = k, Y = 1\} \cup \{S = k, Y = 0\} \\ &= \{S = k\} \cap (\{Y = 1\} \cup \{Y = 0\}) \\ &= \{S = k\} \cap \Omega \\ &= \{S = k\} \end{aligned}$$

and thus, $\mathbb{P}(\hat{Y} = 1 \mid S = k, Y = 1) = \mathbb{P}(\hat{Y} = 1 \mid S = k)$ for all $k = 1, \dots, K$.

Secondly, we also define $\min_k := \min_{k \in \{1, \dots, K\}} \mathbb{P}(\hat{Y} = 1 \mid S = k)$. Then, we easily find that $\text{DI} = \frac{\min_k}{\max_k}$ and, together with the first observation, $\text{EO} = \max_k - \min_k$ holds (cf. definition 2.7 and 2.15).

Now the rest follows by simple equivalent transformations: In setting 1, we find that

$$\frac{\min_k}{\max_k} \geq 1 - \epsilon \iff \min_k \geq (1 - \epsilon) \max_k \iff \max_k - \min_k \leq \epsilon \max_k \quad (3.5)$$

holds. In setting 2, we obtain

$$\max_k - \min_k \leq \tilde{\epsilon} \iff 1 - \frac{\min_k}{\max_k} \leq \frac{\tilde{\epsilon}}{\max_k} \iff \frac{\min_k}{\max_k} \geq 1 - \frac{\tilde{\epsilon}}{\max_k}. \quad (3.6)$$

□

Corollary 3.2. *Given the setting of lemma 3.1,*

1. *$\text{EO} = \tilde{\text{EO}}$ for $\tilde{\text{EO}} := (1 - \text{DI}) \cdot \max_k$ and*

2. *$\text{DI} = \tilde{\text{DI}}$ for $\tilde{\text{DI}} := 1 - \frac{\text{EO}}{\max_k}$ holds.*

Proof. This is a direct consequence of lemma 3.1, where we choose $\epsilon := 1 - \text{DI}$ in setting 1 and $\tilde{\epsilon} := \text{EO}$ in setting 2, and where we can work with equalities instead of estimations in equation (3.5) and (3.6), respectively. □

3.3 Application Domain and Data Set

After having introduced an appropriate definition of a sensitive feature and protected groups in WDSs in the previous subsection 3.2, in order to test whether the concrete instantiations of leakage detection methods presented in subsection 3.1.2 are fair in this sense, we need to generate suitable data based on given WDS structures.

The WDSs considered are *Hanoi* (cf. [33, 41]) and *L-Town* (cf. [39, 40]) displayed in figure 3.2 and 3.3, respectively. While Hanoi consists of 32 nodes, among which three are provided with sensors, and 34 links, L-Town displays a more realistic WDS consisting of 785 nodes, among which 33 are provided with sensors, and 909 links. The latter is constructed in a way to mimic a true WDS while satisfying security defaults.

Pressure measurement simulation For security reasons, only a limited number of real-world data sets based on such systems are available. Therefore, to evaluate methods such as the H-method presented in subsection 3.1.2, data has to be simulated.

For Hanoi, we generate pressure measurements with a time window of $\delta = 10\text{min.}$ using the atmtoolbox (cf. [37]). The pressure is simulated at the sensor nodes displayed in figure 3.2 and for different leakage scenarios, which differ in the leakage location and size. As the WDS is relatively small, we are able to simulate a leakage at each node in the system and for three different diameters $d \in \{5, 10, 15\}\text{cm.}$ In total, the data set is balanced with respect to the label, i.e., the fact whether ($\{Y = 1\}$) or whether not ($\{Y = 0\}$) a leakage present at the time of the considered sample.




Figure 3.2: The Hanoi WDS. The Hanoi WDS, its sensor nodes (IDs 3, 10 and 25) and the protected groups, each highlighted in another color (group 1 on the left side in light shade, group 2 in the middle in dark shade, group 3 on the right side in middle shade). The sensor nodes are colored in the same color of the protected group they belong to and highlighted with a grey circle.

For L-Town, we generate pressure measurements with a time window of $\delta = 5\text{min.}$ as used in the work of Ashraf et al. [6]. The pressure is simulated at the sensor nodes displayed in figure 3.3 and for different leakage scenarios. Due to the larger system size, we are only able to simulate a leakage at some nodes in the system and for three different diameters $d \in \{1.9, 2.3, 2.7\}\text{cm}^{11}$.

¹¹ The different leakage sizes for the two WDSs Hanoi and L-Town can be explained based on some physical background onto which we can not comment in full detail. However, roughly speaking, the size of the chosen leakage diameters is depending on the water supply and demand dynamics of a WDS. For Hanoi, the number of consumers are small, however, water demands are high since the water source




Figure 3.3: The L-Town WDS. The L-Town WDS, its sensor nodes and the protected groups, each highlighted in another color (group 1, also called area C, on the top left in middle shade; group 2, also called area A, in light shade; group 3, also called area B, on the bottom in the middle in dark shade). The sensor nodes are colored in the same color of the protected group they belong to and highlighted with a grey circle.

Pressure residual computation Consecutively, in order to obtain the pressure residuals required for the H- or other method(s), virtual sensor predictions have to be generated (cf. figure 3.1). For Hanoi, we train and use linear virtual sensors with a preprocessing hyperparameter of $T_r = 2$ such as Artelt et al. [5] do. For L-Town, we train and use GCN virtual sensors, for which no hyperparameters are required (cf. subsection 3.1.2).

Protected groups Finally, the protected groups as introduced in subsection 3.2 are displayed in figure 3.2 and 3.3 as well. Here, we work with $K = 3$ different groups for both the Hanoi and the L-Town WDS.

3.4 Experimental Results and Analysis: Residual-Based Ensemble Leakage Detection Does Not Obey Fairness

In table 3.1, the results of the H-method presented in subsection 3.1.2 are shown. The hyperparameter $\Theta \in \mathcal{X} = \mathbb{R}_+^d$ is chosen manually per diameter d such that the test accuracy is close to maximal. On the one hand, we see that independent on the WDS and the virtual sensors used, the method in general performs better in terms of accuracy ACC the larger the leakage size is, as larger leakages are associated with larger pressure drops. Moreover, the method is capable of detecting even small leakages with high(er) accuracy in larger (and therefore, more realistic) WDSs (cf. footnote 11 for details). On the other hand, and more importantly, we see that independent on the WDS and the virtual sensors used, the method is unfair in terms of disparate impact score DI, where a value of 0.8 or larger is desirable (cf. [43]), and equal opportunity score EO. However, the experimental evaluation confirms the mathematical findings of corollary 3.2 by comparing the column of the disparate impact score calculated according to definition 2.7 (DI) to the one according to corollary 3.2.2 (\tilde{DI}), and the column of the equal opportunity score calculated according to definition 2.15 (EO) to the one according to corollary 3.2.1 (\tilde{EO}).

has sufficient water pressure. Hence, larger leakage diameters (proportional to demands) are required to simulate a significant leak. For L-Town, the number of consumers is much larger that have smaller individual demands since the water pressure at the reservoirs is the same as Hanoi. Hence, smaller leakage diameters are sufficient to simulate significant leakages.

d	ACC	\max_k	\min_k	DI	EO	$\tilde{\text{DI}}$	$\tilde{\text{EO}}$
5	0.6223	0.8468	0.4880	0.5763	0.3558	0.5763	0.3588
10	0.7998	0.9983	0.6372	0.6383	0.3611	0.6383	0.3611
15	0.8837	1.0000	0.6402	0.6402	0.3598	0.6402	0.3598

(a) Hanoi WDS and linear virtual sensors.

d	ACC	\max_k	\min_k	DI	EO	$\tilde{\text{DI}}$	$\tilde{\text{EO}}$
1.9	0.7034	0.8935	0.4828	0.5404	0.4107	0.5404	0.4107
2.3	0.8346	1.0000	0.6652	0.6652	0.3348	0.6652	0.3348
2.7	0.8476	1.0000	0.4254	0.4254	0.5746	0.4254	0.5746

(b) L-Town WDS and GCN virtual sensors.

Table 3.1: Results of the H-method. Results of the H-method with \max_k and \min_k according to (the proof of) lemma 3.1. Moreover, the disparate impact and equal opportunity score DI and EO as well as $\tilde{\text{DI}}$ and $\tilde{\text{EO}}$ according to definition 2.7, 2.15, corollary 3.2.2 and 3.2.1, respectively.

This also justifies that in our setting, the usage of one of the two scores is sufficient. Therefore, from now on, we mostly work with the disparate impact score DI only.

4 Fairness-Enhancing Leakage Detection in Water Distribution Networks

Motivated by the result that the standard leakage detection method presented in subsection 3.1.2 does not satisfy the notions of fairness, as another main contribution of this work, we modify this H-method to enhance fairness as introduced in subsection 2.3. The main idea is based on the fact that in the H-method the only models trained are the virtual sensors f_j^r for all $j = 1, \dots, d$ (cf. subsection 3.1.2). However, given these virtual sensors and resulting residuals $r(t_i) \in \mathcal{X} = \mathbb{R}_+^d$, as well as labels $y(t_i) \in \mathcal{Y} = \{0, 1\}$ for times $t_i \in \mathbb{R}$, we can turn the choice of the hyperparameter $\Theta := (\theta_j)_{j=1, \dots, d} \in \mathcal{X}$ of the ensemble classifier $f^c = f^c(\cdot, \Theta)$ (cf. eq. (3.4)) into an optimization problem (OP), where $\Theta \in \mathcal{X}$ now acts as a parameter. The corresponding hypothesis space is $\mathcal{H} := \{f^c : \mathcal{X} \rightarrow \mathcal{Y}, r \mapsto f^c(r, \Theta) \mid \Theta \in \mathcal{X}\}$ (cf. subsection 2.2).

In the following section, we therefore propose (subsection 4.1) and evaluate (subsection 4.2) different, in contrast to the H-method optimization-based, methods that aim at optimizing the parameter $\Theta \in \mathcal{X}$ in order to obtain an optimal ensemble classifier $f^c(\cdot, \Theta_{\text{opt.}}) \in \mathcal{H}$. Optimality hereby depends on the OP at hand: These methods on the one hand are further baselines, where treating the modelling problem as an OP enables us to optimize the result of the H-method itself without fairness considerations. On the other hand, we consider fairness-enhancing methods, where the parameter $\Theta \in \mathcal{X}$ needs to be optimized such that the resulting ensemble classifier is simultaneously as accurate and fair on the given training data as possible.

4.1 Methodology

The following methods define training algorithms to find an optimal ensemble classifier $f^c(\cdot, \Theta_{\text{opt.}}) \in \mathcal{H}$. The scores considered in these algorithms rely on labeled training data

740 $\mathcal{D}^c = \{(r(t_i), y(t_i)) \in \mathcal{X} \times \mathcal{Y} \mid i = 1, \dots, n_c\}^{12}$, which also holds samples based on leaky
 741 states $\{Y = 1\}$ of the WDS. For simplicity, we omit the dependence of all considered
 742 functions on the training data \mathcal{D}^c .

743 4.1.1 Fair Leakage Detection Framework

744 In general, a learning problem such as the training of an optimal ensemble classifier
 745 $f^c(\cdot, \Theta_{\text{opt.}}) \in \mathcal{H}$ can be phrased as an OP, where the objective is to minimize some
 746 suitable loss function $L = L(\Theta) := L(f^c(\cdot, \Theta), \cdot) : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$ over the hypothesis space
 747 \mathcal{H} , or more precisely, with respect to the parameter $\Theta \in \mathcal{X}$, based on its evaluations on
 748 the training data $\mathcal{D}^c \subset \mathcal{X} \times \mathcal{Y}$:

$$749 \quad \left\{ \min_{\Theta \in \mathcal{X}} L(\Theta). \right. \quad (4.7)$$

750 The advantage of redefining the choice of hyperparameters $\Theta \in \mathcal{X}$ (which is what we
 751 do in the H-method) as an OP is that we can now extend this OP by side constraints
 752 $C_k = C_k(\Theta) := C_k(f^c(\cdot, \Theta), \cdot) : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$:

$$753 \quad \left\{ \begin{array}{ll} \min_{\Theta \in \mathcal{X}} & L(\Theta), \\ \text{s.t.} & C_k(\Theta) \geq 0 \quad \forall k = 1, \dots, \hat{K}. \end{array} \right. \quad (4.8)$$

754 A typical way of optimizing a constrained OP is to integrate the side constraints in the
 755 objective in order to apply unconstrained optimization algorithms. This can be done
 756 using a so called *barrier-* or *penalty function* $p : \mathbb{R} \rightarrow [-\infty, \infty]$ (cf. [28]). Using such
 757 functions, the constrained OP (4.8)¹³ can be transformed to

$$758 \quad \left\{ \min_{\Theta \in \mathcal{X}} L(\Theta) + \mu \sum_{k=1}^{\hat{K}} p(C_k(\Theta)). \right. \quad (4.9)$$

759 Hereby, the hyperparameter $\mu \in [0, \infty)$ regulates the importance of the constraints C_k for
 760 all $k = 1, \dots, \hat{K}$ compared to the loss function L .

761 4.1.2 Fair Leakage Detection Instantiations

762 Equation (4.9) gives a general framework on how to train a (fair) leakage detection model
 763 based on the general leakage detection pipeline presented in subsection 3.1.1. While
 764 the H-method presented in subsection 3.1.2 is an instantiation of this pipeline that only
 765 requires the training of the virtual sensors, i.e., the first component of the pipeline, the
 766 following methods also require the training of the leakage detection model itself, i.e., the
 767 third component of the pipeline (cf. figure 3.1).

¹²In practise, we train and test the (ensemble) classifier(s) on unseen data for times $i \geq n_r + 1$. However, for the sake of readability, we choose the indices $i = 1, \dots, n_c$ instead of $i = n_r + 1, \dots, n_c$.

¹³Note that the requirement $\Theta = (\theta_j)_{j=1, \dots, d_c} \in \mathcal{X} = \mathbb{R}_+^d$ actually also contains the constraint $\theta_j \geq 0$ for all $j = 1, \dots, d$. Nevertheless, also the residuals $r(t_i) \in \mathcal{X} = \mathbb{R}_+^d$ are non-negative for all $i = 1, \dots, n$. Therefore, if for any $j \in \{1, \dots, d\}$, $\theta_j < 0$ holds, the ensemble classifier $\hat{Y} = f^c(X, \Theta)$ (cf. eq. (3.4)) is equal to $\hat{Y} \equiv 1$, i.e., it *only* predicts leakages. As our datasets are balanced with respect to the labels, this will lead to an accuracy of approximately 0.5 and to a TPR, but also a FPR of 1. Thus, such choices are no (local) optima of the OPs as they either do not deliver a(n) (locally) optimal loss or as they harm the side constraint(s). In other words, the solution of the OPs will automatically be feasible with respect to the constraint $\theta_j \geq 0$ for all $j = 1, \dots, d$. Therefore, for simplicity, we do not include this constraint as a regularization term in the objective and can optimize Θ over \mathbb{R}^d instead of $\mathcal{X} = \mathbb{R}_+^d$ anyways.

More precisely, the following methods are instantiations of this third component using the framework proposed in the previous subsection 4.1.1. Hereby, the ensemble classifier $f^c(\cdot, \Theta) \in \mathcal{H}$ on which the loss function L and the side constraints C_k for $k = 1, \dots, \hat{K}$ rely on is of the same structure as in the H-method (cf. eq. (3.4)); the resulting optimal models $\hat{Y} = f^c(X, \Theta_{\text{opt.}})$ only differ in their optimal parameter $\Theta_{\text{opt.}} \in \mathcal{X} = \mathbb{R}_+^d$. We obtain such different optimal parameters by choosing different loss functions L , different side constraints C_k for $k = 1, \dots, \hat{K}$ and different algorithmic choices. In the following, we propose possible such choices. The indices (loss index, constraint index, optimization index and barrier or penalty function index) introduced along the way will later be used for the names of the resulting explicit methods as combinations of such choices. A general scheme of this overall idea as an extension of figure 3.1 is displayed in figure 4.1.

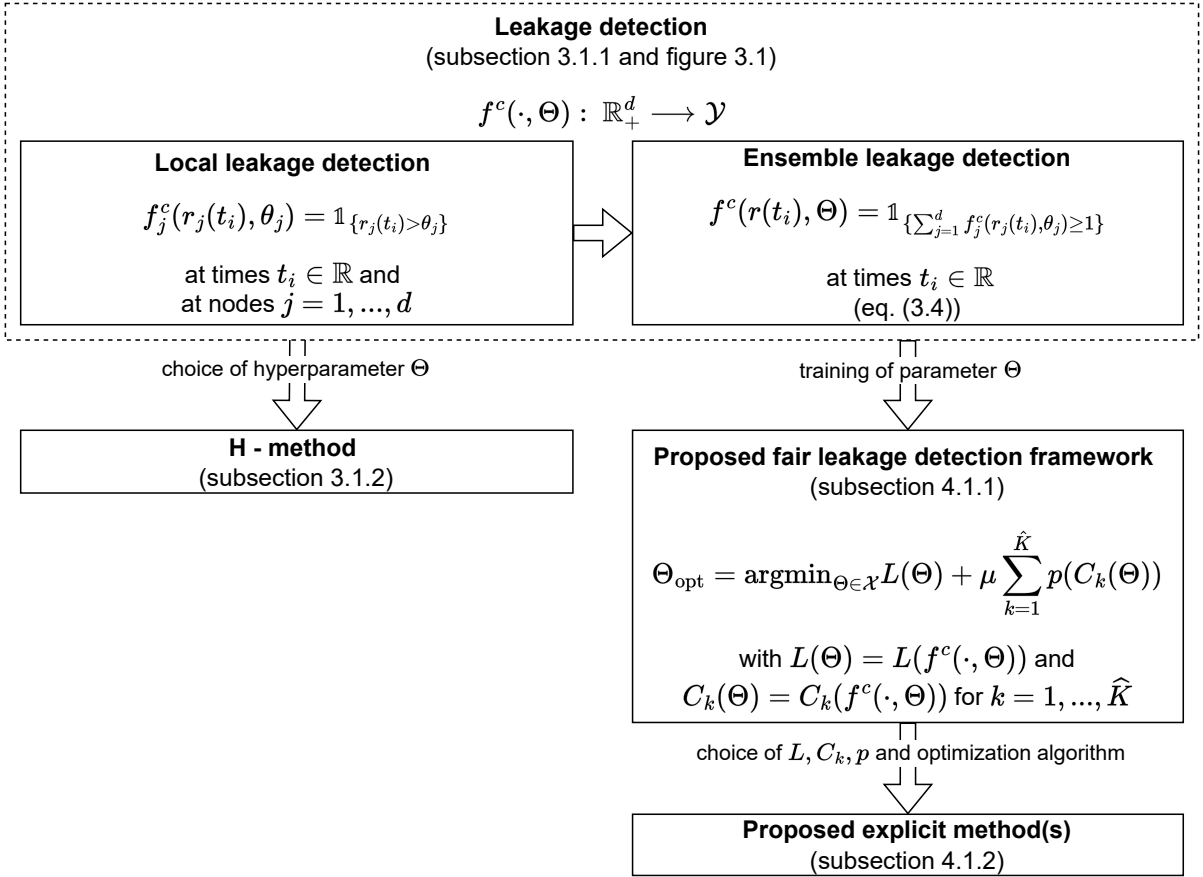


Figure 4.1: Fair leakage detection framework as an extension of figure 3.1.

Optimizing Performance as Baseline Methods By choosing a typical evaluation score as the loss function L and not using any further (fairness) constraints (i.e., $\mu = 0$ or $\hat{K} = 0$), we obtain further baseline methods which output optimized parameters $\Theta_{\text{opt.}} = \operatorname{argmin}_{\Theta \in \mathcal{X}} L(\Theta)$ compared to the H-method and by that, with respect to the performance of the leakage detection model $\hat{Y} = f^c(X, \Theta_{\text{opt.}})$, but not with respect to its fairness. Typical such evaluation scores for a binary classification task are:

- The negative accuracy, i.e. $L(\Theta) = -\text{ACC}(\Theta)$ (loss index “ACC”),
- the negative difference $L(\Theta) = -\text{TPR}(\Theta) + \text{FPR}(\Theta)$ between the TPR and the FPR (loss index “TFPR”).

Optimizing Performance under Fairness Constraints For the following approaches, the loss function L controls the performance while the constraints C_k control the fairness for all $k = 1, \dots, \hat{K}$.

Choice of performance loss functions: When optimizing the performance under fairness constraints, we choose the same loss functions as when optimizing the performance without fairness constraints as introduced in the previous paragraph.

Choice of fairness constraints: As done in our previous work [35], in terms of fairness constraints, we make use of the covariance between the sensitive feature(s) and the prediction of the ensemble classifier. For technical reasons¹⁴, we have to transform the non-binary sensitive feature $S : \{Y = 1\} \rightarrow \{1, \dots, K\}$ to K binary sensitive features $S_k : \Omega \rightarrow \{0, 1\}$, which give answer to the question whether ($\{S_k = 1\}$) or whether not ($\{S_k = 0\}$) a leakage is active in group k for all $k = 1, \dots, K$.¹⁵ Using that $\hat{y}(t_i) = f^c(r(t_i), \Theta)$ holds for all realisations $i = 1, \dots, n_c$, for all binary sensitive features S_k for $k = 1, \dots, K$, the *empirical* covariance between a single sensitive feature S_k and the model $\hat{Y} = f^c(X, \Theta)$ is given by

$$\text{Cov}_{\text{emp.}}(S_k, \hat{Y}) = n_c^{-1} \sum_{i=1}^{n_c} (s_k(t_i) - \bar{s}_k) \cdot f^c(r(t_i), \Theta). \quad (4.10)$$

The usage of the (empirical) covariance as a proxy for fairness is based on the idea that group fairness of a model \hat{Y} , or more precisely, a high disparate impact score on which we focus in this work, relies on the assumption of \hat{Y} being independent of the sensitive feature S (cf. subsection 2.3.2), or in our case, each of the sensitive features S_k for $k = 1, \dots, K$. As independence of two random variables implies their covariance being equal to zero, the latter can be interpreted as a necessary condition for fairness. For more information on this intuition, but also on how our contributions are generalizations of the work of Zafar et al. [43], we refer to our previous work [35].

Motivated by that, we require $\text{Cov}_{\text{emp.}}(S_k, \hat{Y}) \leq c$ and $\text{Cov}_{\text{emp.}}(S_k, \hat{Y}) \geq -c$ to hold, or, equivalently formulated in standard form:

- We require $C_k(\Theta) := c - \text{Cov}_{\text{emp.}}(S_k, \hat{Y}) \geq 0$ and $C_k(\Theta) := c + \text{Cov}_{\text{emp.}}(S_k, \hat{Y}) \geq 0$ to hold for all $k = 1, \dots, K$ (i.e., $\hat{K} = 2K$ in equation (4.8)). Hereby, the hyperparameter $c \in [0, \infty)$ regulates how much the covariance's absolute value is bounded and therefore, the desired fairness (constraint index "COV").

Optimizing Fairness under Performance Constraints For the following approaches, the loss function L controls the fairness while the constraints C_k control the performance for all $k = 1, \dots, \hat{K}$.

Choice of fairness loss functions: As done in our previous work [35], we choose the disparate impact score DI as a loss function. Moreover, as elaborated in the conclusion of our previous work [35] and similar to Zafar et al. [43], we additionally change the role of the empirical covariance by optimizing a fairness proxy similar to the one introduced in equation (4.10) directly. Therefore, taking into account that we have multiple sensitive values, two reasonable loss functions are:

¹⁴As discussed in subsection 3.3, while the model \hat{Y} is defined on the population space Ω , the sensitive feature S is defined on $\{Y = 1\} \subset \Omega$. Moreover, the empirical covariance is only well-defined for variables that are metric scaled or binary nominal scaled.

¹⁵More precisely, $\{S = k\} = \{S_k = 1\}$ holds, with the advantage that the model \hat{Y} and the binary sensitive features S_k for all $k = 1, \dots, K$ are defined on the same domain; on the population space Ω .

- The sum $L(\Theta) := \sum_{k=1}^K |\text{Cov}(S_k, \hat{Y})|$ of absolute values of the empirical covariance between a single sensitive feature S_k and the model $\hat{Y} = f^c(X, \Theta)$ for all $k = 1, \dots, K$ (loss index “Cov”),
- the negative disparate impact score $L(\Theta) := -\text{DI}(\Theta)$ (cf. definition 2.7) (loss index “DI”).

Choice of performance constraints: In terms of performance constraints, we stick to the choice of the accuracy ACC, which is only allowed to differ by some percentage of the optimal accuracy $\text{ACC}_{\text{opt.}}$ obtained when training without fairness constraints (cf. [35, 43]). More precisely, we require $\text{ACC}(\Theta) \geq (1 - \lambda) \text{ACC}_{\text{opt.}}$ or, equivalently formulated in standard form:

- We require $C_1(\Theta) := \text{ACC}(\Theta) - (1 - \lambda) \text{ACC}_{\text{opt.}} \geq 0$ to hold (i.e., $\hat{K} = 1$ in equation (4.8)). Hereby, the hyperparameter $\lambda \in [0, 1]$ regulates how much the accuracy $\text{ACC}(\Theta)$ is allowed to differ from the optimal accuracy $\text{ACC}_{\text{opt.}}$ received, e.g., by another baseline method (constraint index “ACC”).

By that, it indirectly regulates the fairness as well, as the more the accuracy is allowed to differ from the optimal accuracy, the larger the feasible subspace of \mathcal{X} gets and by that, the more the fairness as the loss in the objective can be optimized.

Algorithmic choices Next to the choices of loss function and constraints, the final methods also differ in dependence of what algorithmic choices are made, e.g., what optimization algorithm as well as what barrier or penalty function p is used (cf. eq. 4.9).

One question to answer when choosing an optimization algorithm is whether the considered objective of the OP is (continuously) differentiable. In the setting of ML, the objective clearly depends on the model’s prediction $\hat{Y} = f^c(X, \Theta)$ or more precisely, on $y(t_i) = f^c(r(t_i), \Theta)$ for all $i = 1, \dots, n_c$. However, in view of the ensemble classifier’s definition (cf. eq. (3.4)), f^c is not differentiable with respect to Θ .

Therefore, in dependence on the fact whether we chose a differentiable (db) or non-differentiable (ndb) optimization algorithm, we need to approximate the model:

- If we want to use a gradient-based optimization technique, we make $\hat{Y} = f^c(X, \cdot)$ differentiable by approximating each indicator function $\mathbb{1}_{\{v>0\}}$ by the sigmoid function $\text{sgd}_b(v) = (1 + \exp^{-bv})^{-1}$ with hyperparameter $b \in [0, \infty)$ (optimization index “db”). Replacing the ensemble classifier’s prediction $f^c(r(t_i), \Theta)$ (cf. eq. (3.4)) by

$$\hat{f}^c(r(t_i), \Theta) := \text{sgd}_b \left(\sum_{j=1}^{d_c} \text{sgd}_b(r_j(t_i) - \theta_j) - T \right)$$

for all $i = 1, \dots, n_c$ yields a differentiable approximation of the model \hat{Y} . Hereby, we replace the threshold 1 of the exact ensemble classifier f^c by a hyperparameter $T \in [0, 1]$ to handle the insecurity of the sigmoid function around zero.

- If we want to use a non gradient-based optimization technique, we do not make any changes (optimization index “ndb”).

For more details on that, we refer to our previous work [35]. Depending on what optimization algorithm is used, different (differentiable or non-differentiable) barrier or penalty functions p can be used. In this work, we make use of

- the barrier function $p(t) := -\log(t)$ (barrier function index “log”) and
- the penalty function $p(t) := \max\{0, -t\}$ (penalty function index “max”).

Explicit methods Finally, after having presented all possible choices, we obtain the following explicit methods using the following nomenclature:

loss index+[constraint index–optimization index–barrier/penalty function index].

Each resulting fairness-enhancing method comes with a corresponding baseline method to which it will be compared to in the evaluation:

- the fairness-enhancing TFPR+COV-db-log-method with corresponding baseline TFPR-db-method,
- the fairness-enhancing TFPR+COV-ndb-log- and TFPR+COV-ndb-max-method with corresponding baseline TFPR-ndb-method,
- the fairness-enhancing ACC+COV-db-log-method with corresponding baseline ACC-db-method,
- the fairness-enhancing ACC+COV-ndb-log- and ACC+COV-ndb-max-method with corresponding baseline ACC-ndb-method,
- the fairness-enhancing COV+ACC-ndb-log- and COV+ACC-ndb-max-method also with corresponding baseline ACC-ndb-method and
- the fairness-enhancing DI+ACC-ndb-log- and DI+ACC-ndb-max-method also with corresponding baseline ACC-ndb-method.

The first two notes refer to the fairness-enhancing methods where performance is optimized under fairness constraints and the last four notes refer to the fairness-enhancing methods where fairness is optimized under performance constraints.

4.2 Experimental Results and Analysis

Based on the pressure measurements in the Hanoi WDS and the pressure residuals we obtain from these measurements by making use of the virtual sensors (cf. subsection 3.3), we test all methods introduced in subsection 3.1.2 (H-method) and subsection 4.1.2 in practise. Afterwards, we will test the best performing method on the data associated with the more complex and more realistic L-Town WDS.

Training and testing setup: To test the considered methods, a model \hat{Y} is trained per method and per leakage diameter d on training data (40% of the overall data) and evaluated on test data (60% of the overall data).¹⁶ For the training, the different OPs presented in subsection 4.1 are solved using the BFGS algorithm (cf. [28]) in case of a differentiable OP and the Downhill-Simplex-Search algorithm, also known as the Nelder-Mead algorithm, (cf. [17]) in case of a non-differentiable OP in order to find the optimal parameter $\Theta_{\text{opt.}} \in \mathcal{X}$ of the leakage detection model $\hat{Y} = f^c(X, \Theta_{\text{opt.}})$.

The implementation of all methods and all our results can be found on GitHub¹⁷.

¹⁶Note that since for the fairness evaluation, enough data from all different groups is required, we choose a comparably high percentage of the data for testing.

¹⁷https://github.com/jstrotherm/FairnessInWDSs_extended

4.2.1 Hanoi

Initial parameters Optimization algorithms require an initial start parameter. For the experiments on the Hanoi WDS, we use the hyperparameter $\Theta_{\text{opt.}} \in \mathcal{X}$ found for the H-method (cf. subsection 3.1) as an initial parameter $\Theta_0 \in \mathcal{X}$ for the remaining optimization-based methods (cf. subsection 4.1).

Hyperparameters While the parameters $\Theta_{\text{opt.}} \in \mathcal{X}$ are now outputs of these optimization-based methods, these are subordinate to other hyperparameters. In table 4.1, an overview of these hyperparameters are displayed per method (and if they differ, per diameter). We choose suitable hyperparameters $\mu, b \in [0, \infty)$ and $T \in [0, 1]$ and keep them fixed afterwards. In contrast, the fairness-hyperparameters $c \in [0, \infty)$ or $\lambda \in [0, 1]$, i.e., the hyperparameters that regulate the fairness directly or indirectly, respectively, are changed to obtain different *score combinations* of performance, measured by the accuracy score ACC, and fairness, measured by the disparate impact or equal opportunity score DI or EO, respectively. We do so by starting with a hyperparameter c or λ that causes perfect fairness, i.e., a disparate impact score of 1.0, whenever possible and in- or decrease the hyperparameter by 0.01 until the disparate impact score of the resulting fairness-enhanced model achieves an equal or worse disparate impact score than its corresponding baseline method, respectively (cf. paragraph “Explicit methods” in subsection 4.1.2 or table 4.1 for the corresponding baseline method per fairness-enhancing method).

Method	$c \in [0, \infty)$	$\lambda \in [0, 1]$	μ (d = 5, 10, 15)	b	T
TFPR-db (b)	-	-	-	100	0.8
TFPR+COV-db-log	✓	-	0.10 0.20 0.20	100	0.8
TFPR-ndb (b)	-	-	-	-	-
TFPR+COV-ndb-log	✓	-	0.20 0.25 0.25	-	-
TFPR+COV-ndb-max	✓	-	100	-	-
ACC-db (b)	-	-	-	100	0.8
ACC+COV-db-log	✓	-	0.15 0.05 0.05	100	0.8
ACC-ndb (b)	-	-	-	-	-
ACC+COV-ndb-log	✓	-	0.2 0.3 0.05	-	-
ACC+COV-ndb-max	✓	-	100	-	-
COV+ACC-ndb-log	-	✓	0.01 0.01 0.01	-	-
COV+ACC-ndb-max	-	✓	100	-	-
DI+ACC-ndb-log	-	✓	0.05 0.025 0.04	-	-
DI+ACC-ndb-max	-	✓	100	-	-

Table 4.1: Overview of the used hyperparameters. The “b” indicates baseline methods that aim at optimizing general performance without fairness considerations. For more details on these hyperparameters, see subsection 4.1.2.

Results With these settings in mind, we obtain the following results. As we in total test five baseline methods (the H-method and the ones proposed in subsection 4.1.2) and ten fairness-enhancing methods (cf. subsection 4.1.2), and by that, many methods, we do only present the key findings in this section and further detailed findings regarding the comparison of all these method in appendix B.

Moreover, for a better overview of the results, we divide the ten fairness-enhancing methods into four subcategories: The TFPR-methods including all methods with loss index “TFPR”, and analogously the ACC-methods, the COV-methods and the DI-methods. In some of the results, these methods are represented together with their corresponding baseline methods. Note that two methods from the same subcategory can have different baseline methods as corresponding baseline methods (cf. paragraph “Explicit Methods” in subsection 4.1.2 or table 4.1).

Increasing fairness: In figure 4.2, we see the performance and fairness of some exemplary trained ensemble classifiers measured in accuracy and disparate impact score, respectively. For the fairness-enhancing methods, testing different hyperparameters c or λ causes error bars for these methods. The height of the bars with error bars corresponds to the mean accuracy and disparate impact score achieved by each method over all hyperparameter values tested. The error bars themselves reach from the lowest to the largest score of the two scores considered.




Figure 4.2: Accuracy and disparate impact score per method and leakage diameter in the Hanoi-WDS as well as for different hyperparameters c or λ .

We see that the fairness-enhancing methods on average increase fairness while on average decreasing accuracy compared to their corresponding baseline methods. However, the average increase in fairness is larger than the average decrease in accuracy. For details regarding different diameters d , the score ranges and the other methods, we refer to appendix B. Based on these, one can say that fairness and overall performance are mutually dependent to about the same extent.

In addition to that, figure 4.3 shows the performance and indirectly, also the fairness of some exemplary trained ensemble classifiers measured by the TPR per group. The height of the bars and the range of the error bars behave analogously to figure 4.2.




Figure 4.3: TPR per method, group and leakage diameter in the Hanoi-WDS as well as for different hyperparameters c or λ .

In view of the definition of the equal opportunity score (cf. definition 2.15) and due to the fact that this score is equivalent to the disparate impact score in our domain of application (cf. lemma 3.1 and corollary 3.2), a method is fairer in our context the more similar the TPRs per group are. This is what we observe in figure 4.3 (and figure B.2) when comparing the TPRs among groups for the fairness-enhancing methods to the TPRs among groups for their corresponding baseline methods. Even more, figure 4.3 (and figure B.2) show(s) that the increase in fairness that we observe in figure 4.2 (and figure B.1) on average is not only obtained by decreasing the performance of the (in the corresponding baseline method) best-performing group, but also by increasing the performance of the (in the corresponding baseline method) worst-performing group. By some methods, even all TPRs per groups are increased on average.

The coherence of fairness and overall performance, and non-optimality: While figure 4.2 and 4.3 only hint at the relationship between fairness and overall performance, measured in disparate impact and accuracy score, respectively, a more detailed visualization of how fairness is related to the overall performance of a model can be found in figure 4.4. For each tested hyperparameter c or λ , respectively, depending on what fairness-enhancing method was used, the obtained score combinations, i.e., the accuracy and the disparate impact score, are visualized for some exemplary trained ensemble classifiers.

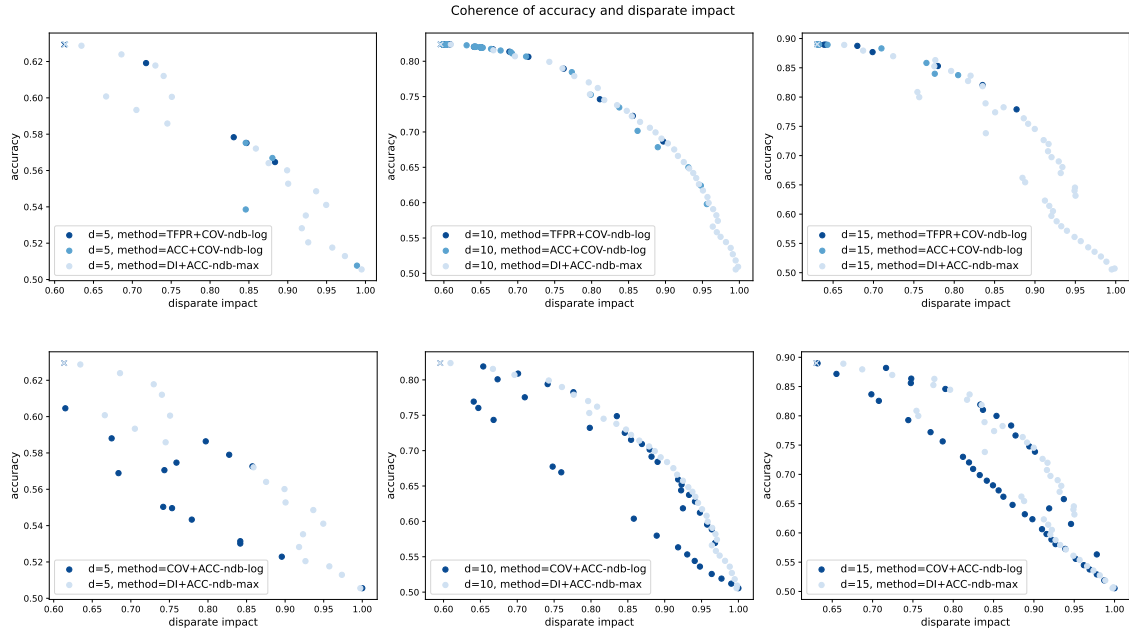


Figure 4.4: Coherence of accuracy and disparate impact score for the different fairness-enhancing methods and different leakage sizes in the Hanoi-WDS, based on different hyperparameters c or λ . The cross data points visualize the accuracy and disparate impact score of the corresponding baseline methods (cf. paragraph “Explicit Methods” in subsection 4.1.2 or table 4.1).

The characteristic curve that can be observed in most of all subimages is the so-called Pareto front, visualizing that the increase in fairness is accompanied by the reduction in accuracy score and vice versa. Note that the non-optimal solutions apart from the Pareto-front in figure 4.4 and also later on, the local jumps recognized in figure 4.5, can be explained by the non-convexity of the objective functions. Because of that, the found

solutions $\Theta_{\text{opt.}} \in \mathcal{X}$ strongly depend on the initialized parameter $\Theta_0 \in \mathcal{X}$ and might not correspond to the global optimum.

Nevertheless, by most fairness-enhancing methods, a desired disparate impact score of about 0.8 can be achieved by a decrease of accuracy by approximately 0.03 - 0.06 points below the optimal accuracy obtained by the corresponding baseline methods (depending on the specific method used). Hereby, both fairness and overall performance can be influenced by the fairness-hyperparameters c or λ , respectively. Deciding which choice of fairness-hyperparameter is optimal and by that, deciding on the trade-off between fairness and overall performance, is a difficult task that depends on the extent of the decisions of the underlying model as well as legal requirements. Regarding legal requirements, by not using the sensitive features for the decision making of the algorithms, the methods presented can satisfy the legal definition of disparate treatment *and* disparate impact (depending on the hyperparameter chosen) simultaneously.

Another observation is that the largest accuracies of the fairness-enhancing methods are usually approximately as good as the accuracy of their corresponding baseline methods while achieving equal or better fairness results. In opposite direction, perfect fairness of $\text{DI} = 1.0$ can be achieved at a cost of the worst possible accuracy of $\text{ACC} = 0.5$. Depending on the method, the jump in disparate impact and accuracy score is rather abrupt or more fine-grained when reaching this extreme of $(\text{DI}, \text{ACC}) = (1.0, 0.5)$: Especially the COV- and the DI-methods relying on the optimization of fairness while constraining the accuracy using the hyperparameter λ allow the latter, because the accuracy constraint is less sensitive than the covariance constraints, controlled by the hyperparameter c .¹⁸

However, also some of the TFPR- and the ACC-methods relying on the optimization of performance while constraining the fairness using the hyperparameter c allow fine-grained variations in both scores. This motivates to investigate into the different methods also within the chosen subcategories. We do so on appendix B.

Here, we find that the DI+ACC-ndb-max-method provides the best results while also providing the benefit of only requiring a few hyperparameters which are easy to choose. This finding makes the DI+ACC-ndb-max-method the best candidate to be tested on a more complicated and by that, more realistic, WDS, as we will do in subsection 4.2.2. However, before we do so, we investigate more on the relation between the performance and fairness scores and the fairness-hyperparameters c and λ .

The influence of the fairness-hyperparameters on fairness and overall performance: In figure 4.5, for the best-performing method of the TFPR-methods and the DI-methods – the results for the ACC-methods look similar to the ones of the TFPR-methods and the results for the COV-methods look similar to the ones of the DI-methods –, we show how the hyperparameters are related to disparate impact and accuracy, but this time, also equal

¹⁸This is because too small choices of c cause possible solutions with penalty or barrier function(s) of infinity. In such case, the trivial solution of only predicting leakages remains left, as in this case, the covariance becomes zero (cf. eq. (4.10) for $f^c(r(t_i), \Theta) = 1$ for all $i = 1, \dots, n$) and by that, the penalty or the barrier function(s) are *not* infinity. Even more, in this case, all TPRs per group and by that, the disparate impact score, are equal to $\text{DI} = 1$. Moreover, in this case, the accuracy score is approximately $\text{ACC} = 0.5$ as the data set used is balanced with respect to the labels (cf. subsection 3.3). Therefore, in such cases, we end up with the trivial combination of $(\text{DI}, \text{ACC}) = (1.0, 0.5)$. In contrast, starting with a hyperparameter λ that causes the trivial solution and then decreasing this hyperparameter to enforce a larger accuracy until the optimal accuracy is achieved ($\lambda = 0$) easily allows optimizing fairness without harming the accuracy constraint as it is measured in units of the optimal accuracy, that - as proven by the corresponding baseline method - exists.

opportunity score. Each of the three scores is plotted against the used hyperparameter of the corresponding fairness-enhancing method tested.

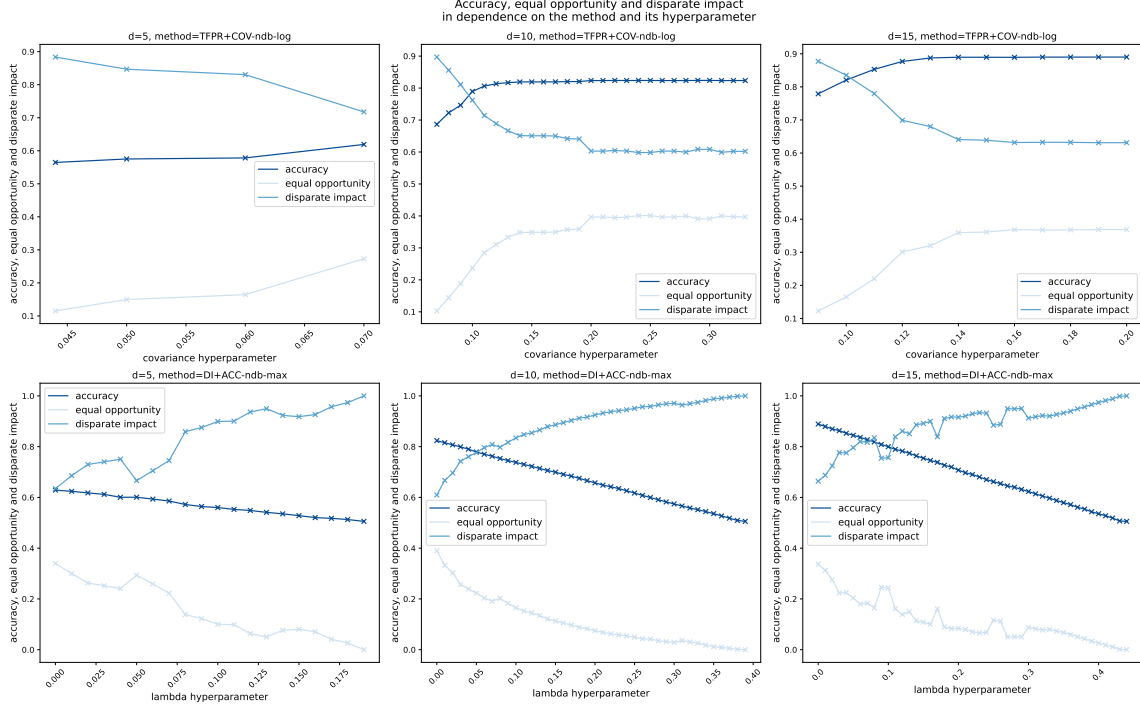


Figure 4.5: Coherence of accuracy, disparate impact, equal opportunity and the training hyperparameter for different fairness-enhancing methods and different leakage sizes in the Hanoi-WDS.

For the TFPR+COV-ndb-log-method (and the ACC+COV-ndb-log-method), the decrease of the hyperparameter c is accompanied by the improvement of the fairness measures as well as the decrease of the performance measure. This can be explained as follows: A high empirical covariance between a sensitive feature and the prediction of the ensemble classifier means that the relative number of positive predictions within the related group differs significantly from the relative number of positive predictions within a group with small covariance. Thus, the more the covariance is constrained by the hyperparameter c , the less such extreme differences in the relative number of positive predictions across groups occur, leading to a better fairness score. In the case of disparate impact, therefore, a (better) higher score at the expense of a (worse) lower overall performance - compared to the overall performance that occurs in the unconstrained case or for a looser constraint, that is a larger bound c , - appears. In the case of equal opportunity, however, a (better) lower score at the expense of a (worse) lower overall performance appears.

In contrast, for the DI+ACC-ndb-max-method (and the COV+ACC-ndb-log-method), the increase of the hyperparameter λ is accompanied by the improvement of the fairness measures as well as the decrease of the performance measure due to the fact that a higher hyperparameter λ allows a larger deviation of the optimal accuracy score. Thus, the feasible search space is extended and a worse accuracy is penalized less or not at all, so that the fairness score in the objective can be optimized to a larger extend.

Equivalence of disparate impact and equal opportunity: Moreover, especially to mention is the observation of our theoretical results from lemma 3.1 and corollary 3.2 in practise: For

the coherence of equal opportunity score and the hyperparameters, the results in figure 4.5 equal the ones for disparate impact score in the same figure, but reflected along the horizontal axis through the point (0,0.5). This proves the equivalence of both fairness measures as theoretically proven in lemma 3.1 and corollary 3.2. Nevertheless, note that this is an application specific result and does not hold in general.

Finally, as another new contribution compared to our previous work [35], we will test the best-performing DI+ACC-ndb-max-method on a more complex and by that realistic WDS, L-Town, using the more powerful GCN-virtual sensors incorporated into the leakage detection method.

4.2.2 L-Town

Initial parameters While the dimension of the search space $\mathcal{X} = \mathbb{R}_+^d$ is equal to $d = 3$ (with d the number of sensors) in Hanoi, it extends to $d = 33$ in L-Town (cf. subsection 3.3). By that, chances are high that the graph of the objective function that needs to be optimized in each of the presented optimization-based methods (cf. subsection 4.1) gets more complex and exhibits more saddle points and local minima. This intuition turns out to be true in practise, where the choice of the initial start parameters $\Theta_0 \in \mathcal{X}$ are crucial to the success of the methods tested. Therefore, for the experiments on the L-Town WDS, we use the hyperparameter $\Theta_{\text{opt.}} \in \mathcal{X}$ found for the H-method (cf. subsection 3.1) only as an initial parameter $\Theta_0 \in \mathcal{X}$ for the ACC-ndb-method, which is the corresponding baseline method for the DI+ACC-ndb-max-method (cf. paragraph “Explicit Methods” in subsection 4.1.2 or table 4.1) that turned out to work best in the previous subsection 4.2.1. Using the same initial parameter for the DI+ACC-ndb-max-method itself did not provide optimal results (– the pareto-front obtained here did not end up in the score combination of the corresponding baseline method). Therefore, consecutively, we use the hyperparameter $\Theta_{\text{opt.}} \in \mathcal{X}$ found by the ACC-ndb-method as an initial parameter $\Theta_0 \in \mathcal{X}$ for the DI+ACC-ndb-max-method.

Hyperparameters In view of table 4.1, the ACC-ndb-method does not require the choice of any hyperparameters. For the DI+ACC-ndb-max-method, we vary the fairness-hyperparameter $\lambda \in [0, 1]$ and also choose $\mu = 100$ as discussed in subsection 4.2.1.

Results Similar to figure 4.4 for Hanoi, figure 4.6 shows the relation between the fairness and the overall performance of the trained model applied to L-Town.

The observations for L-Town are similarly well compared to those on Hanoi. Although while it first seems that there are less score combinations apart from, or more precisely, *below*, the pareto-front compared to the results of the same method applied to Hanoi, some score combinations *above* the seemingly optimal pareto-front may give raise to the existence of an even better pareto-front, which is not observed completely due to non-convexity of the OP.

Nevertheless, a desired disparate impact score of about 0.8 can be achieved by a decrease of accuracy by approximately 0.1 points for $d = 1.9$ and 0.01 points for $d = 2.3$ below the optimal accuracy obtained. For $d = 2.7$, the leakages are already almost detected perfectly and fair by the corresponding baseline ACC-ndb-method. Anyways, the fairness-enhancing DI-ACC-ndb-max-method is better by approximately 0.015 points in disparate impact score with barely no loss in accuracy.




Figure 4.6: Coherence of accuracy and disparate impact score for the DI+ACC-ndb-max-method and different leakage sizes in the L-Town-WDS. The cross data points visualize the accuracy and disparate impact score of the corresponding baseline, the ACC-method.

Finally, similar to figure 4.5 for Hanoi, figure 4.7 shows how the hyperparameters are related to accuracy, disparate impact and equal opportunity score in the setting of L-Town. The results go hand in hand with the observations found for Hanoi, and also the equivalence between the two fairness scores can be observed again.

Additionally, we see by the position of the accuracy curves and the slope of the fairness curves that on the one hand, the better the model performs in general, measured by the accuracy score, the fairer the model is initially, and on the other hand, the harder it is to make the model even fairer.




Figure 4.7: Coherence of accuracy, disparate impact, equal opportunity and the training hyperparameter for the DI+ACC-ndb-max-method and different leakage sizes in the L-Town-WDS.

5 Conclusion

In this work, we introduced the notion of group fairness in an application domain of high social and ethical relevance, namely in the field of water distribution systems (WDSs). This required the generalization of common group fairness definitions to a single or possibly multiple non-binary sensitive feature(s). To do so, we gave a detailed introduction on the concept of group fairness based on the mathematical concept of independence, derived these generalized definitions from this concept and proved that they coincide with common group fairness definitions in the case of a binary sensitive feature and a binary classification task. We then investigated on the fairness issue in the area of leakage de-

tection within WDSs. We showed that standard approaches are not fair in the context of different groups related to the locality within the network. As a remedy, we presented multiple methods that increase fairness of the leakage detection model with respect to the introduced fairness notion while satisfying the legal notions of disparate treatment and disparate impact simultaneously. We tested these methods not only on the Hanoi WDS, but also on the more complex and by that more realistic L-Town WDS. We empirically demonstrated that fairness and overall performance of the model are interdependent and the use of hyperparameters provides the ability to trade off fairness and overall performance. However, this trade off lies in the responsibility of the policy maker. Nevertheless, the notion of fairness within the water domain is still in its beginning and further work on other cases of application within this domain is crucial.

6 Funding Statement

We gratefully acknowledge funding from the European Research Council (ERC) under the ERC Synergy Grant Water-Futures (Grant agreement No. 951424).

List of Abbreviations

ACC accuracy.

AI artificial intelligence.

DI disparate impact.

DP demographic parity.

EO equal opportunity.

EOs equalized odds.

ERC European Research Council.

EU European Union.

FPR false positive rate.

GCN graph convolutional network.

iff if and only if.

ML machine learning.

OP optimization problem.

TNR true negative rate.

TPR true positive rate.

WDS water distribution system.

References

- [1] Alekh Agarwal, Alina Beygelzimer, Miroslav Dudík, John Langford, and Hanna Wallach. A Reductions Approach to Fair Classification. In *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pages 60–69. PMLR, 2018.
- [2] Alekh Agarwal, Miroslav Dudík, and Zhiwei Steven Wu. Fair Regression: Quantitative Definitions and Reduction-based Algorithms. In *Proceedings of the 36th International Conference on Machine Learning*, *Proceedings of Machine Learning Research*, pages 120–129. PMLR, 2019.
- [3] Sina Aghaei, Mohammad Javad Azizi, and Phebe Vayanos. Learning Optimal and Fair Decision Trees for Non-Discriminative Decision-Making. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 1418–1426, 2019.
- [4] J. Angwin, J. Larson, S. Mattu, and L. Lauren Kirchner. Machine Bias - There’s software used across the country to predict future criminals. And it’s biased against blacks. *ProPublica*, 2016. Online available at <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>; accessed at 21/07/2023.
- [5] André Artelt, Stelios Vrachimis, Demetrios Eliades, Marios Polycarpou, and Barbara Hammer. One Explanation to Rule them All – Ensemble Consistent Explanations. *arXiv preprint arXiv:2205.08974*, 2022.
- [6] Inaam Ashraf, Luca Hermes, André Artelt, and Barbara Hammer. Spatial Graph Convolution Neural Networks for Water Distribution Systems. In *Proceedings of the 21th International Symposium on Intelligent Data Analysis (IDA)*, *Lecture Notes in Computer Science*, pages 29–41. Springer, 2023.
- [7] Solon Barocas, Moritz Hardt, and Arvind Narayanan. *Fairness and Machine Learning: Limitations and Opportunities*. fairmlbook.org, 2019. <http://www.fairmlbook.org>.
- [8] Heinz Bauer. *Probability Theory*, volume 23. Walter de Gruyter, 1996.
- [9] Richard Berk, Hoda Heidari, Shahin Jabbari, Matthew Joseph, Michael Kearns, Jamie Morgenstern, Seth Neel, and Aaron Roth. A Convex Framework for Fair Regression. *arXiv preprint arXiv:1706.02409*, 2017.
- [10] Dan Biddle. *Adverse Impact and Test Validation: A Practitioner’s Guide to Valid and Defensible Employment Testing*. Gower Publishing, Aldershot, UK, 2006.
- [11] Toon Calders, Asim Karim, Faisal Kamiran, Wasif Ali, and Xiangliang Zhang. Controlling Attribute Effect in Linear Regression. In *Proceedings of the 13th IEEE International Conference on Data Mining*, pages 71–80. IEEE, 2013.
- [12] Alessandro Castelnovo, Riccardo Crupi, Greta Greco, Daniele Regoli, Ilaria Giuseppina Penco, and Andrea Claudio Cosentini. A clarification of the nuances in the fairness metrics landscape. *Scientific Reports*, 12(1):4209, 2022.

- [13] Sam Corbett-Davies, Emma Pierson, Avi Feller, Sharad Goel, and Aziz Huq. Algorithmic Decision Making and the Cost of Fairness. In *Proceedings of the 23rd International Conference on Knowledge Discovery and Data Mining (SIGKDD)*, pages 797–806. ACM Digital Library, 2017.
- [14] Cynthia Dwork, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Richard Zemel. Fairness Through Awareness. In *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference (ITCS)*, pages 214–226. ACM Digital Library, 2012.
- [15] European Commission, Directorate-General for Communications Networks, Content and Technology. *Ethics guidelines for trustworthy AI*. Publications Office, 2019. Online available at <https://data.europa.eu/doi/10.2759/346720>; accessed at 21/07/2023.
- [16] Michael Feldman, Sorelle A Friedler, John Moeller, Carlos Scheidegger, and Suresh Venkatasubramanian. Certifying and Removing Disparate Impact. In *Proceedings of the 21th International Conference on Knowledge Discovery and Data Mining (SIGKDD)*, pages 259–268. ACM Digital Library, 2015.
- [17] Fuchang Gao and Lixing Han. Implementing the Nelder-Mead simplex algorithm with adaptive parameters. *Computational Optimization and Applications*, 51(1):259–277, 2012.
- [18] Guancheng Guo, Xipeng Yu, Shuming Liu, Ziqing Ma, Yipeng Wu, Xiyan Xu, Xiaoting Wang, Kate Smith, and Xue Wu. Leakage Detection in Water Distribution Systems Based on Time-Frequency Convolutional Neural Network. *Journal of Water Resources Planning and Management*, 147(2), 2021.
- [19] Moritz Hardt, Eric Price, and Nati Srebro. Equality of Opportunity in Supervised Learning. In *Proceedings of the 29th International Conference on Advances in Neural Information Processing Systems (NeurIPS)*, 2016.
- [20] Rolf Isermann. *Fault-Diagnosis Systems – An Introduction from Fault Detection to Fault Tolerance*. Springer, Berlin, Heidelberg, Germany, 2006.
- [21] Faisal Kamiran and Toon Calders. Classifying without Discriminating. In *Proceedings of the 2nd International Conference on Computer, Control and Communication (IC4)*, pages 1–6. IEEE, 2009.
- [22] Faisal Kamiran and Toon Calders. Classification with No Discrimination by Preferential Sampling. In *Proceedings of the 19th Machine Learning Conference of Belgium and The Netherlands*, volume 1, 2010.
- [23] Achim Klenke. *Wahrscheinlichkeitstheorie*, volume 1. Springer, 2006.
- [24] Junpei Komiyama, Akiko Takeda, Junya Honda, and Hajime Shimao. Nonconvex Optimization for Regression with Fairness Constraints. In *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pages 2737–2746. PMLR, 2018.
- [25] Zhirong Li, Jiaying Wang, Hexiang Yan, Shuping Li, Tao Tao, and Kunlun Xin. Fast Detection and Localization of Multiple Leaks in Water Distribution Network Jointly

Driven by Simulation and Machine Learning. *Journal of Water Resources Planning and Management*, 148(9), 2022.

[26] Ninareh Mehrabi, Fred Morstatter, Nripsuta Saxena, Kristina Lerman, and Aram Galstyan. A Survey on Bias and Fairness in Machine Learning. *ACM Computing Surveys (CSUR)*, 54(6):1–35, 2021.

[27] Harikrishna Narasimhan, Andrew Cotter, Maya Gupta, and Serena Wang. Pairwise Fairness for Ranking and Regression. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 5248–5255, 2020.

[28] Jorge Nocedal and Stephen J. Wright. *Numerical Optimization*, volume 02. Springer, New York, NY, 2006.

[29] Dana Pessach and Erez Shmueli. A Review on Fairness in Machine Learning. *ACM Computing Surveys (CSUR)*, 55(3):1–44, 2022.

[30] Geoff Pleiss, Manish Raghavan, Felix Wu, Jon Kleinberg, and Kilian Q Weinberger. On Fairness and Calibration. In *Proceedings of the 30th International Conference on Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

[31] Luis Romero-Ben, Débora Alves, Joaquim Blesa, Gabriela Cembrano, Vicenç Puig, and Eric Duviella. Leak Localization in Water Distribution Networks using Data-Driven and Model-Based Approaches. *Journal of Water Resources Planning and Management*, 148(5), 2022.

[32] Boris Ruf and Marcin Detyniecki. Towards the Right Kind of Fairness in AI. *arXiv preprint arXiv:2102.08453*, 2021.

[33] Ildeberto Santos-Ruiz, Francisco-Ronay López-Estrada, Vicenç Puig, Guillermo Valencia-Palomo, and Héctor-Ricardo Hernández. Pressure Sensor Placement for Leak Localization in Water Distribution Networks Using Information Theory. *Sensors*, 22(2):443, 2022.

[34] David B Steffebauer, Jochen Deuerlein, Denis Gilbert, Edo Abraham, and Olivier Piller. Pressure-Leak Duality for Leak Detection and Localization in Water Distribution Systems. *Journal of Water Resources Planning and Management*, 148(3), 2022.

[35] Janine Strotherm and Barbara Hammer. Fairness-Enhancing Ensemble Classification in Water Distribution Networks. In *Proceedings of the 17th International Work-Conference on Artificial Neural Networks (IWANN)*, volume 14134 of *Lecture Notes in Computer Science*, pages 119–133. Springer, 2023.

[36] Janine Strotherm, Alissa Müller, Barbara Hammer, and Benjamin Paaßen. Fairness in KI-Systemen. *arXiv preprint arXiv:2307.08486*, 2023.

[37] Jonas Vaquet. Automation Toolbox for Machine learning in water Networks, 2023. Online available at <https://pypi.org/project/atmn/>.

[38] Michael Veale and Frederik Zuiderveen Borgesius. Demystifying the Draft EU Artificial Intelligence Act – Analysing the good, the bad, and the unclear elements of the proposed approach. *Computer Law Review International*, 22(4):97–112, 2021.

- 1250 [39] Stelios G Vrachimis, Demetrios G Eliades, Riccardo Taormina, Zoran Kapelan, Avi
1251 Ostfeld, Shuming Liu, Marios Kyriakou, Pavlos Pavlou, Mengning Qiu, and Marios M
1252 Polycarpou. Battle of the Leakage Detection and Isolation Methods. *Journal of Water
1253 Resources Planning and Management*, 148(12), 2022.
- 1254 [40] Stelios G. Vrachimis, Demetrios G. Eliades, Riccardo Taormina, Avi Ostfeld, Zoran
1255 Kapelan, Shuming Liu, Marios Kyriakou, Pavlos Pavlou, Mengning Qiu, and Mar-
1256 ios M. Polycarpou. BattLeDIM: Battle of the Leakage Detection and Isolation Meth-
1257 ods. In *Proceedings of the 2nd International Joint Conference on Water Distribution
1258 Systems Analysis & Computing and Control for the Water Industry (WDSA/CCWI)*,
1259 pages 1–6, 2020.
- 1260 [41] Stelios G. Vrachimis, Marios S. Kyriakou, Demetrios G. Eliades, and Marios M.
1261 Polycarpou. LeakDB: A Benchmark Dataset for Leakage Diagnosis in Water Distri-
1262 bution Networks. In *Proceedings of the 1st International Joint Conference on Wa-
1263 ter Distribution Systems Analysis & Computing and Control for the Water Industry
1264 (WDSA/CCWI)*, 2018.
- 1265 [42] Muhammad Bilal Zafar, Isabel Valera, Manuel Gomez Gomez Rodriguez, and Kr-
1266 ishna P Gummadi. Fairness Beyond Disparate Treatment & Disparate Impact:
1267 Learning Classification Without Disparate Mistreatment. In *Proceedings of the 26th
1268 International Conference on World Wide Web (WWW)*, pages 1171–1180. ACM Dig-
1269 ital Librabry, 2017.
- 1270 [43] Muhammad Bilal Zafar, Isabel Valera, Manuel Gomez Rogriguez, and Krishna P
1271 Gummadi. Fairness Constraints: Mechanisms for Fair Classification. In *Proceedings
1272 of the 20th International Conference on Artificial Intelligence and Statistics (AIS-
1273 TATS)*, volume 54 of *Proceedings of Machine Learning Research*, pages 962–970.
1274 PMLR, 2017.

A Additional Proofs

A.1 Sigma-Fields

Lemma A.1 (σ -field induced by a set system).

Let Ω be a set and $\mathcal{E} \subset P(\Omega)$ some set system of Ω . Then the set system

$$\sigma(\mathcal{E}) := \bigcap_{\substack{\mathcal{F} \text{ is } \sigma\text{-field on } \Omega, \\ \mathcal{E} \subset \mathcal{F}}} \mathcal{F}$$

defines the smallest σ -field on Ω that holds \mathcal{E} , called the σ -field generated by \mathcal{E} .

Proof. This is trivial. \square

Remark A.2. For the set system $\mathcal{E} \subset P(\Omega)$ of Ω , $\mathcal{E} \subset \sigma(\mathcal{E})$ holds. If the set system $\mathcal{E} \subset P(\Omega)$ is a σ -field on Ω already, $\sigma(\mathcal{E}) = \mathcal{E}$ holds: All σ -fields \mathcal{F} on Ω considered in the intersection satisfy $\mathcal{E} \subset \mathcal{F}$. Therefore, $\mathcal{E} \subset \bigcap_{\mathcal{F} \text{ is } \sigma\text{-field on } \Omega, \mathcal{E} \subset \mathcal{F}} \mathcal{F}$ holds. If \mathcal{E} is a σ -field already, \mathcal{E} is one of these σ -fields on Ω that satisfy $\mathcal{E} \subset \mathcal{E}$. Therefore, in this case, $\bigcap_{\mathcal{F} \text{ is } \sigma\text{-field on } \Omega, \mathcal{E} \subset \mathcal{F}} \mathcal{F} \subset \mathcal{E}$ also holds.

Remark A.3. If for two set systems $\mathcal{E}_1, \mathcal{E}_2 \subset P(\Omega)$ of Ω , $\mathcal{E}_1 \subset \mathcal{E}_2$ holds, $\sigma(\mathcal{E}_1) \subset \sigma(\mathcal{E}_2)$ holds: All σ -fields that hold \mathcal{E}_2 also hold \mathcal{E}_1 . Therefore, the intersection over all σ -fields that hold \mathcal{E}_1 creates an equal or smaller set system compared to the intersection over all σ -fields that hold \mathcal{E}_2 .

Lemma A.4 (σ -field induced by a function).

Let $X : \Omega \rightarrow \mathcal{X}$ be a function with σ -field $\mathcal{F}_{\mathcal{X}}$ on \mathcal{X} . Then the set system $\sigma(X) := X^{-1}(\mathcal{F}_{\mathcal{X}}) = \{X^{-1}(A) \subset \Omega \mid A \in \mathcal{F}_{\mathcal{X}}\}$ of Ω defines a σ -field on Ω , called the σ -field generated by X .

Proof. (i) $\Omega \in \sigma(X)$: By definition of a σ -field on \mathcal{X} , $\mathcal{X} \in \mathcal{F}_{\mathcal{X}}$ holds. Thus, by definition of $\sigma(X)$, $\Omega = X^{-1}(\mathcal{X}) \in \sigma(X)$ holds.

(ii) If $B \in \sigma(X)$, $B^C \in \sigma(X)$, too: If $B \in \sigma(X)$, by definition of $\sigma(X)$, there exists an $A \in \mathcal{F}_{\mathcal{X}}$, such that $B = X^{-1}(A)$ holds. By definition of a σ -field, $A^C \in \mathcal{F}_{\mathcal{X}}$ holds. Thus, by definition of $\sigma(X)$, $B^C = (X^{-1}(A))^C = X^{-1}(A^C) \in \sigma(X)$ holds.

(iii) If $B_n \in \sigma(X)$ for all $n \in \mathbb{N}$, $\bigcup_{n \in \mathbb{N}} B_n \in \sigma(X)$, too: If $B_n \in \sigma(X)$, by definition of $\sigma(X)$, there exists an $A_n \in \mathcal{F}_{\mathcal{X}}$, such that $B_n = X^{-1}(A_n)$ holds for all $n \in \mathbb{N}$. By definition of a σ -field, $\bigcup_{n \in \mathbb{N}} A_n \in \mathcal{F}_{\mathcal{X}}$ holds. Thus, by definition of $\sigma(X)$, $\bigcup_{n \in \mathbb{N}} B_n = \bigcup_{n \in \mathbb{N}} X^{-1}(A_n) = X^{-1}(\bigcup_{n \in \mathbb{N}} A_n) \in \sigma(X)$ holds. \square

Lemma A.5 (σ -field induced by a set system and a function).

Let $X : \Omega \rightarrow \mathcal{X}$ be a function and $\mathcal{E} \subset P(\mathcal{X})$ some set system of \mathcal{X} . Then $X^{-1}(\sigma(\mathcal{E})) = \sigma(X^{-1}(\mathcal{E}))$ holds.

Proof. (i) $\sigma(X^{-1}(\mathcal{E})) \subset X^{-1}(\sigma(\mathcal{E}))$: By remark A.3, $X^{-1}(\mathcal{E}) \subset X^{-1}(\sigma(\mathcal{E}))$ implies $\sigma(X^{-1}(\mathcal{E})) \subset \sigma(X^{-1}(\sigma(\mathcal{E})))$. By lemma A.4 for $\mathcal{F}_{\mathcal{X}} = \sigma(\mathcal{E})$, $X^{-1}(\sigma(\mathcal{E}))$ is a σ -field on Ω . Thus, by remark A.2, $\sigma(X^{-1}(\mathcal{E})) \subset \sigma(X^{-1}(\sigma(\mathcal{E}))) = X^{-1}(\sigma(\mathcal{E}))$ holds.

(ii) $X^{-1}(\sigma(\mathcal{E})) \subset \sigma(X^{-1}(\mathcal{E}))$: By definition of $X^{-1}(\sigma(\mathcal{E}))$, we need to show that for all $A \in \sigma(\mathcal{E})$, $X^{-1}(A) \in \sigma(X^{-1}(\mathcal{E}))$ holds. We do so by using the principle of good sets:

Let $\mathcal{G} := \{A \subset \mathcal{X} \mid X^{-1}(A) \in \sigma(X^{-1}(\mathcal{E}))\}$. The goal is to show that $\sigma(\mathcal{E}) \subset \mathcal{G}$ holds.

(ii.1) $\mathcal{E} \subset \mathcal{G}$: If $A \in \mathcal{E}$, by definition of $X^{-1}(\mathcal{E})$ and remark A.2, $X^{-1}(A) \in X^{-1}(\mathcal{E}) \subset \sigma(X^{-1}(\mathcal{E}))$ holds. Thus, $A \in \mathcal{G}$ holds.

(ii.2) \mathcal{G} is a σ -field on \mathcal{X} :

(ii.2.i) $\mathcal{X} \in \mathcal{G}$: By definition of a σ -field on Ω , $X^{-1}(\mathcal{X}) = \Omega \in \sigma(X^{-1}(\mathcal{E}))$ holds. Thus, by definition of \mathcal{G} , $\mathcal{X} \in \mathcal{G}$ holds.

(ii.2.ii) If $A \in \mathcal{G}$, $A^C \in \mathcal{G}$, too: If $A \in \mathcal{G}$, by definition of \mathcal{G} , $X^{-1}(A) \in \sigma(X^{-1}(\mathcal{E}))$ holds. By definition of a σ -field, $X^{-1}(A^C) = (X^{-1}(A))^C \in \sigma(X^{-1}(\mathcal{E}))$ holds. Thus, by definition of \mathcal{G} , $A^C \in \mathcal{G}$ holds.

(ii.2.iii) If $A_n \in \mathcal{G}$ for all $n \in \mathbb{N}$, $\cup_{n \in \mathbb{N}} A_n \in \mathcal{G}$, too: If $A_n \in \mathcal{G}$, by definition of \mathcal{G} , $X^{-1}(A_n) \in \sigma(X^{-1}(\mathcal{E}))$ holds for all $n \in \mathbb{N}$. By definition of a σ -field, $X^{-1}(\cup_{n \in \mathbb{N}} A_n) = \cup_{n \in \mathbb{N}} X^{-1}(A_n) \in \sigma(X^{-1}(\mathcal{E}))$ holds. Thus, by definition of \mathcal{G} , $\cup_{n \in \mathbb{N}} A_n \in \mathcal{G}$ holds.

Finally, as $\mathcal{E} \subset \mathcal{G}$ and \mathcal{G} is a σ -field on \mathcal{X} , by remark A.3 and A.2, $\sigma(\mathcal{E}) \subset \sigma(\mathcal{G}) = \mathcal{G}$ holds, which was the goal to show. \square

Remark A.6. If in lemma A.4, the generator of $\mathcal{F}_{\mathcal{X}}$ is known, i.e., if $\mathcal{F}_{\mathcal{X}} = \sigma(\mathcal{E}_{\mathcal{X}})$ holds, using the notation from lemma A.5 with $\mathcal{E} = \mathcal{E}_{\mathcal{X}}$, we obtain

$$\sigma(X) := X^{-1}(\mathcal{F}_{\mathcal{X}}) = X^{-1}(\sigma(\mathcal{E}_{\mathcal{X}})) = \sigma(X^{-1}(\mathcal{E}_{\mathcal{X}})),$$

i.e., the σ -field generated by the function X equals the σ -field generated by the pre-image of the generator $\mathcal{E}_{\mathcal{X}}$ of the σ -field $\mathcal{F}_{\mathcal{X}}$.

A.2 Independence of Two Random Variables

Lemma A.7 (Independence of two random variables, version 1).

X and Y are independent with respect to \mathbb{P} iff

$$\mathbb{P}(X \in A, Y \in B) = \mathbb{P}(X \in A) \cdot \mathbb{P}(Y \in B) \quad (\text{A.11})$$

holds for all $A \in \mathcal{F}_{\mathcal{X}}, B \in \mathcal{F}_{\mathcal{Y}}$.

Proof. By definition of $\sigma(X)$ and $\sigma(Y)$ (cf. definition 2.1) and the definition of independence of two families of events (cf. [8]), $\sigma(X)$ and $\sigma(Y)$ are independent iff

$$\mathbb{P}(X^{-1}(A) \cap Y^{-1}(B)) = \mathbb{P}(X^{-1}(A)) \cdot \mathbb{P}(Y^{-1}(B))$$

holds for all $A \in \mathcal{F}_{\mathcal{X}}, B \in \mathcal{F}_{\mathcal{Y}}$.¹⁹ Using that

$$\begin{aligned} & X^{-1}(A) \cap Y^{-1}(B) \\ &= \{\omega \in \Omega \mid X(\omega) \in A\} \cap \{\omega \in \Omega \mid Y(\omega) \in B\} \\ &= \{\omega \in \Omega \mid X(\omega) \in A, Y(\omega) \in B\} \end{aligned}$$

holds and that $\{X \in A, Y \in B\}$ is just a short form for the latter set, together with analog arguments for $\{X \in A\}$ and $\{Y \in B\}$, we obtain equation (A.11). \square

¹⁹By definition of a random variable, X is \mathcal{F} - $\mathcal{F}_{\mathcal{X}}$ - and Y is \mathcal{F} - $\mathcal{F}_{\mathcal{Y}}$ -measurable, i.e., for all $A \in \mathcal{F}_{\mathcal{X}}$, $X^{-1}(A) \in \mathcal{F}$ and for all $B \in \mathcal{F}_{\mathcal{Y}}$, $Y^{-1}(B) \in \mathcal{F}$ holds. Therefore, the considered probabilities are well defined (\mathbb{P} is a function defined on \mathcal{F}).

Lemma A.8 (Independence of two random variables, version 2).

Assume that $\mathcal{F}_X = \sigma(\mathcal{E}_X)$ and $\mathcal{F}_Y = \sigma(\mathcal{E}_Y)$ holds and that the set systems \mathcal{E}_X and \mathcal{E}_Y , also called generators, are \cap -stable²⁰. Then X and Y are independent with respect to \mathbb{P} iff

$$\mathbb{P}(X \in A, Y \in B) = \mathbb{P}(X \in A) \cdot \mathbb{P}(Y \in B) \quad (\text{A.12})$$

holds for all $A \in \mathcal{E}_X, B \in \mathcal{E}_Y$.

Proof. By definition 2.1, X and Y are independent iff $\sigma(X)$ and $\sigma(Y)$ are independent. By definition of these σ -fields (cf. definition 2.1), then X and Y are independent iff $X^{-1}(\sigma(\mathcal{E}_X))$ and $Y^{-1}(\sigma(\mathcal{E}_Y))$ are independent. By lemma A.5, $X^{-1}(\sigma(\mathcal{E}_X)) = \sigma(X^{-1}(\mathcal{E}_X))$ and $Y^{-1}(\sigma(\mathcal{E}_Y)) = \sigma(Y^{-1}(\mathcal{E}_Y))$ holds. Therefore, X and Y are independent iff $\sigma(X^{-1}(\mathcal{E}_X))$ and $\sigma(Y^{-1}(\mathcal{E}_Y))$ are independent.

For the latter case, using that a probability measure \mathbb{P} is uniquely determined by an \cap -stable generator of the σ -field it is defined on (cf. [23], lemma 1.42), it suffices to test equation (A.11) on $X^{-1}(\mathcal{E}_X)$ and $Y^{-1}(\mathcal{E}_Y)$, respectively, as these are intersection stable if \mathcal{E}_X and \mathcal{E}_Y are. \square

Remark A.9 (\cap -stable generators are enough). If \cap -stable generators \mathcal{E}_X and \mathcal{E}_Y of the σ -fields \mathcal{F}_X and \mathcal{F}_Y , respectively, are known, the following lemmata A.10 and A.12 are replaceable by a version where the σ -fields \mathcal{F}_X and \mathcal{F}_Y are replaced by their generators \mathcal{E}_X and \mathcal{E}_Y , such as we did in lemma A.8 based on lemma A.7.

Lemma A.10 (Independence of two random variables, version 3).

X and Y are independent with respect to \mathbb{P} iff

$$\mathbb{P}(X \in A) = \mathbb{P}(X \in A \mid Y \in B) \quad (\text{A.13})$$

holds for all $A \in \mathcal{F}_X$ and $B \in \mathcal{F}_Y$ for which $\mathbb{P}(Y \in B) > 0$ holds.

Proof. For $A \in \mathcal{F}_X$ and $B \in \mathcal{F}_Y$ for which $\mathbb{P}(Y \in B) > 0$ holds, by definition of conditional probabilities (cf. [8]), $\mathbb{P}(X \in A, Y \in B) = \mathbb{P}(X \in A \mid Y \in B) \cdot \mathbb{P}(Y \in B)$ holds. Comparing equation (A.12) and (A.13) yields both implications, noting that for the case $\mathbb{P}(Y \in B) = 0$, equation (A.12) is trivially fulfilled (cf. remark A.11). \square

Remark A.11. Conditional probabilities of the kind $\mathbb{P}(X \in A \mid Y \in B)$ are only well-defined for $B \in \mathcal{F}_Y$ for which $\mathbb{P}(Y \in B) > 0$ holds. However, equation (A.11) is also satisfied for $B \in \mathcal{F}_Y$ for which $\mathbb{P}(Y \in B) = 0$ holds, because due to rules of (probability) measures, $0 \leq \mathbb{P}(X \in A, Y \in B) \leq \mathbb{P}(Y \in B) = 0$ and therefore, $\mathbb{P}(X \in A, Y \in B) = 0$ holds.

While the results of the previous lemmata are well-known observations, we need a slightly different characterization of independence of random variables than usual to link it to group fairness notions in ML.

Lemma A.12 (Independence of two random variables, version 4).

X and Y are independent with respect to \mathbb{P} iff

$$\mathbb{P}(X \in A \mid Y \in B_1) = \mathbb{P}(X \in A \mid Y \in B_2) \quad (\text{A.14})$$

holds for all $A \in \mathcal{F}_X$ and $B_1, B_2 \in \mathcal{F}_Y$ for which $\mathbb{P}(Y \in B_1), \mathbb{P}(Y \in B_2) > 0$ holds.

²⁰A set system \mathcal{E} is called \cap -stable iff for any two sets $A_1, A_2 \in \mathcal{E}$, also $A_1 \cap A_2 \in \mathcal{E}$ holds.

Proof. If X and Y are independent with respect to \mathbb{P} , equation (A.14) clearly holds for all $A \in \mathcal{F}_X$ and $B_1, B_2 \in \mathcal{F}_Y$ for which $\mathbb{P}(Y \in B_1), \mathbb{P}(Y \in B_2) > 0$ holds by lemma A.10.

Vice versa, if equation (A.14) holds for all $A \in \mathcal{F}_X$ and $B_1, B_2 \in \mathcal{F}_Y$ for which $\mathbb{P}(Y \in B_1), \mathbb{P}(Y \in B_2) > 0$ holds, we prove that X and Y are independent with respect to \mathbb{P} using lemma A.10 as well. To do so, let $A \in \mathcal{F}_X$ and $B \in \mathcal{F}_Y$ for which $\mathbb{P}(Y \in B) > 0$ holds.

Case 1: If $\mathbb{P}(Y \in B) = 1$ holds, by (1) rules of (probability) measures and (2) the definition of conditional probabilities, we obtain

$$0 \stackrel{(1)}{\leq} \mathbb{P}(X \in A, Y \in B^C) \stackrel{(1)}{\leq} \mathbb{P}(Y \in B^C) \stackrel{(1)}{=} 1 - \mathbb{P}(Y \in B) = 0, \text{ and therefore,}$$

$$\begin{aligned} \mathbb{P}(X \in A) &\stackrel{(1)}{=} \mathbb{P}(X \in A, Y \in B) + \underbrace{\mathbb{P}(X \in A, Y \in B^C)}_{=0} \\ &\stackrel{(2)}{=} \mathbb{P}(X \in A \mid Y \in B) \cdot \underbrace{\mathbb{P}(Y \in B)}_{=1} \\ &= \mathbb{P}(X \in A \mid Y \in B). \end{aligned}$$

Case 2: If $0 < \mathbb{P}(Y \in B) < 1$ holds, by definition of a σ -field, $B^C \in \mathcal{F}_Y$ and by the assumption of $\mathbb{P}(Y \in B) < 1$, $\mathbb{P}(Y \in B^C) > 0$ holds. Then, the following conditional probabilities are well-defined and we can use equation (A.14): By (1) rules of (probability) measures, (2) the definition of conditional probabilities and (3) this equation (A.14), we obtain

$$\begin{aligned} \mathbb{P}(X \in A) &\stackrel{(1)}{=} \mathbb{P}(X \in A, Y \in B) + \mathbb{P}(X \in A, Y \in B^C) \\ &\stackrel{(2)}{=} \mathbb{P}(X \in A \mid Y \in B) \cdot \mathbb{P}(Y \in B) + \mathbb{P}(X \in A \mid Y \in B^C) \cdot \mathbb{P}(Y \in B^C) \\ &\stackrel{(3)}{=} (\mathbb{P}(Y \in B) + \mathbb{P}(Y \in B^C)) \cdot \mathbb{P}(X \in A \mid Y \in B) \\ &\stackrel{(1)}{=} \mathbb{P}(X \in A \mid Y \in B). \end{aligned}$$

□

B Additional Experimental Results and Analysis

In this section, we present further detailed findings regarding the comparison of all the methods introduced in subsection 4.1.2.

Increasing fairness: In figure B.1, we see the extension of figure 4.2 for the missing trained ensemble classifiers.

We see that for all fairness-enhancing methods and all leakage sizes, the fairness-enhancing methods on average increase fairness while on average decreasing accuracy by the same or even a smaller amount compared to their corresponding baseline methods. For $d = 5$, all fairness-enhancing methods allow a large range of fairness improvement at cost of a small range of accuracy, which is only due to the relatively poor accuracy of the leakage detection in this scenario in general. Also for $d = 15$ and the TFPR- and the ACC-methods with log-barrier function, the range of fairness is larger than the range of accuracy. However, a perfect fairness score of disparate impact being equal to one can not be achieved by these methods. Such result usually comes along with a low accuracy and would increase the

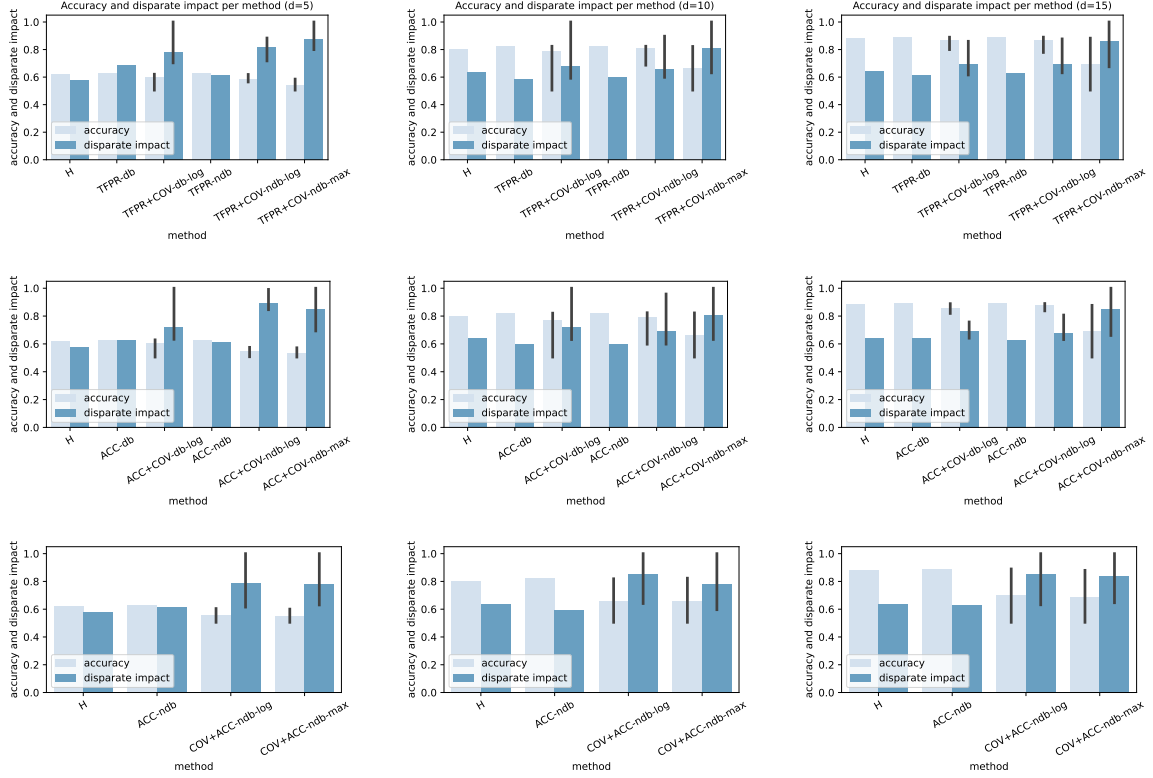


Figure B.1: Accuracy and disparate impact score per method and leakage diameter in the Hanoi-WDS as well as for different hyperparameters c or λ .

accuracy range, as we will see later. For the other scenarios, the ranges of fairness and accuracy are mostly similarly large. Thus, overall, one can say that fairness and overall performance are mutually dependent to about the same extent.

In figure B.2, we see the extension of figure 4.3 for the missing trained ensemble classifiers. The results do not differ significantly compared to figure 4.3.

The coherence of fairness and overall performance: In figure B.3, we see the extension of figure 4.4 for the missing trained ensemble classifier.

Based on our discussions in subsection 3.4, we investigate into the different methods also within the chosen subcategories:

For the TFPR-methods, the methods using the log-barrier yield better results in the sense that their pareto fronts lie above the one using the max-penalty. However, the max-penalty method allows the most fine-grained score combinations, followed by the non-differentiable log-barrier method. As the latter nevertheless allows a disparate impact score larger than 0.8 with a better or similar accuracy score compared to the other TFPR-methods, the TFPR+COV-ndb-log-method is the best performing method among all TFPR-methods.

For the ACC-methods, we observe similar results except that for $d = 10$, the ACC+COV-ndb-log-method even allows the most fine-grained score combinations.

For the COV-methods, the log-barrier method also performs better than the max-penalty method in terms of the position of the pareto-front, but also exhibits some score combinations apart from the pareto-front due to non-convexity of the OP. The non-convexity problem mostly appears for $d = 5$. Nevertheless, as both methods allow fine-grained

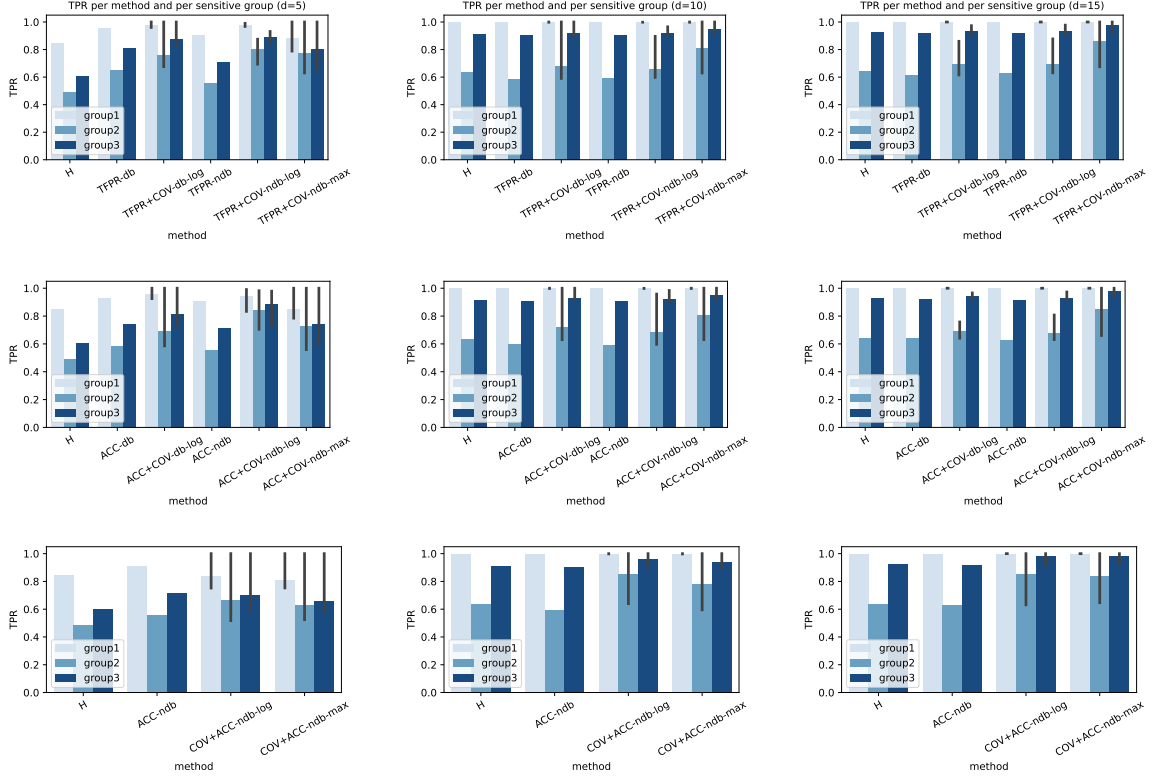


Figure B.2: TPR per method, group and leakage diameter in the Hanoi-WDS as well as for different hyperparameters c or λ .

score combinations, the COV+ACC-ndb-log-method outperforms the COV+ACC-ndb-max-method.

In contrast, for the DI-methods, both log-barrier and max-penalty method provide similarly good pareto-fronts. However, as the max-penalty method allows a little less score combinations apart from the pareto-front and a little more fine-grained score combinations, the DI+ACC-ndb-max-method outperforms the DI+ACC-ndb-log-method.

Also overall, the DI+ACC-ndb-max-method provides the best results: The pareto-front has one of the best shapes (coming closest to the optimal score combination of $(DI, ACC) = (1, 1)$) and is finest-grained while having only a few combinations apart from its curve.

By that, although for all other subcategories, the log-barrier delivers better results, the best method uses the max-penalty, yielding no clear winner of both of them (cf. paragraph “Algorithmic choices” in subsection 4.1.2). Nevertheless, a large advantage of the max-penalty is the easy choice of the hyperparameter μ : While in this case, μ can be any large number (for us, $\mu = 100$ works), for the log-barrier, the choice of μ requires more finetuning (cf. table 4.1).

In contrast, rather more obvious is the result that the non-differentiable methods tend to cause better results compared to the differentiable methods, yielding that the error we make when approximating the ensemble classifier is not compensated by the power of the differentiable optimization algorithm (cf. paragraph “Algorithmic choices” in subsection 4.1.2). Another advantage of the non-differentiable methods are that there are less hyperparameters to choose from (cf. table 4.1).

Overall, the result that the DI+ACC-ndb-max-method provides the best results aligns well with the fact that for this method, the choice of hyperparameters is easiest: While the

hyperparameter μ is easy to choose here as discussed above, also the hyperparameter λ allows a better control of the fairness compared to the hyperparameter c as also elaborated above.

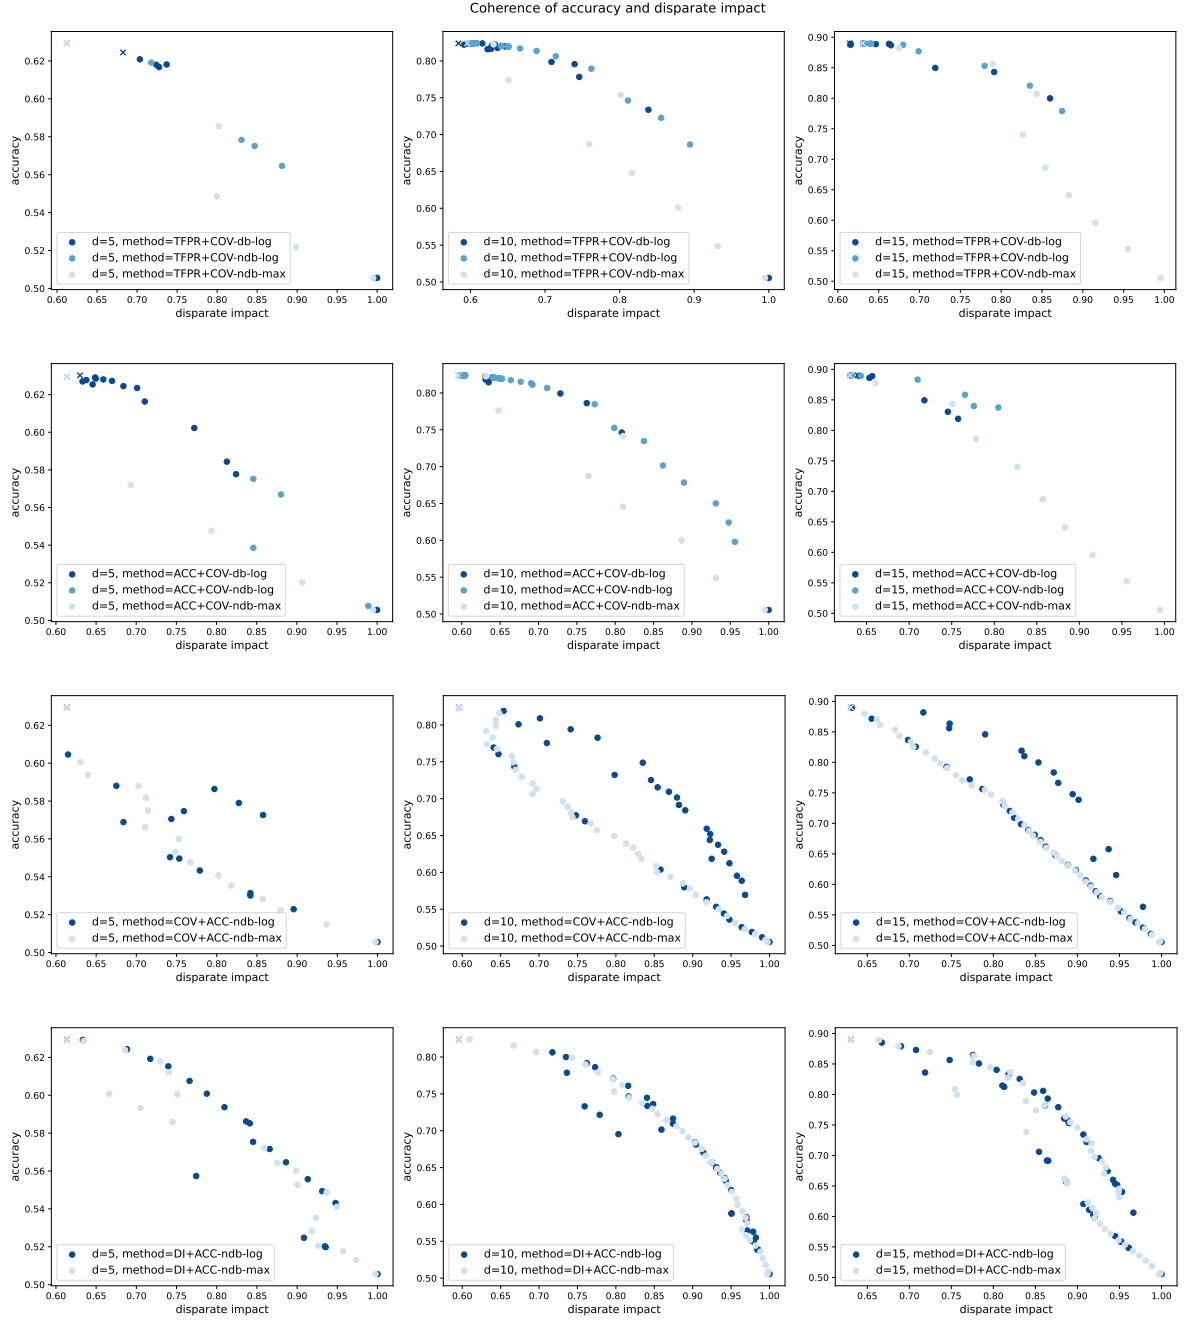


Figure B.3: Coherence of accuracy and disparate impact score for the different fairness-enhancing methods and different leakage sizes in the Hanoi-WDS, based on different hyperparameters c or λ . The cross data points visualize the accuracy and disparate impact score of the corresponding baselines methods (cf. paragraph “Explicit Methods” in subsection 4.1.2 or table 4.1).